Monash University

Electrical Engineering

Final Year Project (ECE4093)

Final Report

Author:
James Anastasiou
ID: 23438940

Supervisor:
Dr. David Boland

September 27, 2016

Contents

1   Significant Contribution

- Designed and implemented a Static Analysis tool for identifying potential parallelism in sequential C code.

- Designed and implemented an AST Matcher Combinator header only library.

- Designed and implemented a CPU and GPU benchmarking suite.

- Designed and implemented a Templated Typeclass interface for parallel primitives in CUDA C++.

2 Introduction

2.1 Purpose and Scope

This document is the complete final report of my final year project. It is the complete documentation of the entire project and as such it includes:

- Project Description

- Literature Review

- Introduction to Parallel Computings

- Overview of Clang Integration

- Detailed Description of Project Components

- Case Study

- Evaluation of Goals and Requirements

- Project Limitations and Possible Extension

- Appendix

3   Project Description

My final year project aimed to create a set of static code/compiler analytics tools to help determine which algorithms within a code-base may be easily parallelized. The specific intention of it is to provide users with a report describing which parts of their code-base may see a potential speed-up from redeveloping them as CUDA (GPU) [1] kernels. In order to achieve this both benchmarking and theoretical analysis was required, as well as static analysis of the C description of the algorithm itself. My final project utilised a combined strategy of static AST analysis, via a hook into the Clang Libtooling, as well as integration with a benchmarking suite I created that utilised real high performance libraries for both Host and Device code. The use case of this software is for developers working primarily in modelling and mathematically intensive areas, as these tend to provide significant opportunity to provide improvement. The expectation of potential users is that they will note that the report alerts them to potential speedups, and then use the performance metrics to determine whether or not to hire a specialised engineer to redesign the relevant components of their system.

3.1   Compiler Analytics

Given a C description of an algorithm, a report is to be generated, highlighting areas within the code that may see a potential speed-up based on matching to known GPU performant patterns. In order to achieve this, a static analysis methodology was invoked. Direct source code analysis is incredibly difficult, and suffers from a number of drawbacks, including the difficulty to parse C and (notoriously difficult) C++. Considering the time constraints involved it was decided that the analysis tool would hook into the Clang Libtooling library, which provides access to the Clang Abstract Syntax Tree, giving a semi-syntax invariant canvas from which to identify relevant parallel patterns.

Given the difficulty of setting up a generic build environment for the Clang tooling, I decided to fork a project named OCLint [2], a C, C++ and Obj-C static analysis tool that I had worked with earlier. This tool provides a light framework around the Clang tooling, however most importantly it provides sophisticated build scripts which work on a variety of operating systems and distributions.

3.1.1 OCLint Modification

Forking the OCLint project, which is licensed under a modified BSD license has saved significant time and effort from being wasted developing a generic build system around the Clang tooling. The OCLint software provides a direct method for interacting with the Clang AST by exposing the Clang Libtooling headers. This increased flexibility has in turn allowed for more time to be spent developing methods to identify parallel patterns within the generated AST. All changes to the OCLint software are unrelated to its original intention and design, and as such no pull requests were lodged, and no modifications I have written have moved upstream. As I have substantially re-engineered and re-purposed the OCLint software I have elected to give it a new name, the C Algorithm Parallelisation Analyser (CAPA).

3.2 GPU Benchmarking

In order to best provide theoretical performance improvements of algorithms within a codebase, an analysis of the current hardware available is significantly important. As such rather than just provide purely theoretical numbers, part of this project involved developing a simple set of GPU benchmarks which seek to show performance metrics for the identified patterns within the code analyser. This in effect means that reports generated by the analytics tool may contain specific information pertaining to the hardware available on the current build and test system. In order to achieve the best outcome, CUDA was decided on as the framework for development.

3.2.1 Benchmarks

GPUs are exceptionally good at high throughput calculations, one particular example is SIMD, meaning *Same Instruction Multiple Data*. The performance of GPUs and the algorithms they are particularly useful for is well under continual research, however general problem classes that GPUs are able to solve efficiently are well understood. These problem sets include algorithms that can be described by any of the following:

- Map

- Fold/Reduce

- Scan/Prefix Network

- Matrix Operations

The actual speed improvements derived from redeveloping serial code to take advantage of the massively parallel compute power of a GPU differs between each of these operations, however many serial algorithms have equivalent or more performant alternate parallel implementations. As such this project involves developing a small set of benchmarks for GPUs that determine their performance in each of these categories. In order to satisfy time constraints and recognise real world concerns, I elected to use existing optimised libraries for the individual components of the benchmarking. I relied on the Eigen Library [3] for host side matrix operations. On the device side I relied on a combination of the thrust template library [4] in combination with CuBLAS [5]

3.2.2 CUDA

CUDA is Nvidia's proprietary library and toolchain for developing parallel software. There are 2 main frameworks in the GPU programming space, CUDA and OpenCL. Whilst OpenCL is a FOSS platform, the development tools are severely lacking in comparison to the CUDA toolkit, and as such it was an easy decision to follow through with the CUDA. This however has limited the performance metrics to only comparisons involving CUDA enabled graphics cards. This is not too great a concern however, as the benchmarking module is highly extensible and as a result can easily integrate with a variety of backends, including OpenCL or OpenMP.

4 Literature Review

4.1 Introduction

Optimisation and computational efficiency are two pillars of good program design, much research has been undertaken in the search of improving performance and extracting hardware maximum efficiency. Although the current literature covers an extensive range of research, this review seeks to focus primarily on the topics of automatic vectorisation, alternative hardware (GPU/FPGA) performance in parallel contexts, and finally the utilisation of Static Analysis and Profiling to assist in the process of identifying potential optimisation in the massively parallel computation paradigm. Individually these are all large topics of research, and as such this review will be focusing primarily on computational performance, rather than power consumption or algorithm efficiency. The purpose of this review is to provide a contextualisation around my final year project, in order to identify potential challenges in the problem space, as elucidated by prior research.

4.2 Automatic Vectorisation

Automatic vectorisation is a tool employed by many compiler designers in order to generate ASM which utilises specialised hardware level vector instructions. Same Instruction Multiple Data (SIMD) instructions seek to process multiple elements of a dataset simultaneously, utilising multiple processing units in order to achieve data level parallelism. Roger Espasa and Mateo Valero [6] explore the potential benefits of Data-Level Parallelism by investigating computer architectures to utilise both Instruction Level and Data Level parallel constructs. Within their research they identify and develop an architecture to utilise SIMD instructions to leverage the performance benefits, clearly demonstrating the performance improvement this parallel strategy provides. These techniques have since been further developed by Compiler designers, with the LLVM/Clang team employing two forms of automatic vectorisation within their optimising C compiler [7]. The Clang team focused on developing Loop Level Vectorisation and Super Word Level Vectorisation in order to leverage parallelisation in the target architecture. Their automatic loop vectoriser is capable of providing an increase in processing speed of 3 times, when tuned specifically for the Intel Core-i7 AVX instruction

set. This is a clear demonstration of the benefits already being seen by optimising compilers manipulating sequential programs into those which leverage the power of parallel computation. The LLVM optimiser however has some flaws which are identified by Yulei Sui, Xiaokang Fan, Hao Zhou, and Jingling Xue who developed Loop-oriented array and field sensitive Pointer Analysis (LPA) in order to combat the sub-optimal performance of the LLVM auto-vectoriser [8]. The authors identified alias analysis as being a potential cause of the LLVM compiler missing optimisation opportunities. They created an analysis framework around both flow-insensitive and flow-sensitive pointer construct. By hooking into the LLVMs partial SSA form they exploited the reduced semantic complexity to algorithmically generate superior memory patterns, and by extension developed a superior loop vectoriser, with performance up to 10% better than the original LLVM. This research again demonstrates the performance opportunities created by parallel computation, however they are all fundamentally limited by the CPUs capacity to perform SIMD operations. Most CPU architectures have at most 8 complex computing cores, limiting the maximum potential throughput, however GPU architectures have the capacity for massively parallel computation. The typical design inherently leverage SIMD concepts containing thousands of simple computational cores with high memory bandwidth. Thus whilst automatic vectorisation has substantial performance potential, most current literature is focused on continuing to improve computational efficiency of host bound programs, rather than looking towards exploiting the massively parallel computational power of modern GPUs.

4.3 Parallelism

General Purpose GPU programming seeks to exploit the parallel performance characteristics of the GPU architecture; identification and development of algorithms which leverage this design pattern can provide substantial computational speedups. The authors of [9] explore the fundamentals of GPGPU programming, and the CUDA architecture. Inherent within the CUDA architecture is hybrid CPU-GPU programming to produce the most efficient solution. In the example described the authors extract maximum parallelism from Matrix operations, and leave complex control flow to the CPU, thus developing a solution which maximises the performance of the individual components of the Hybrid Host-Device

model. The proposed solution utilises parallel primitives such as the Reduction, which exploits parallelism by utilising a summation tree. This requires that the data and binary operator form a semigroup (The set of data must be closed under an associative binary operator). In their example the set is of integers, and the operator is addition, which holds these properties. The concept of parallel primitive operations is further expanded by [10] which provides terminology to describe parallel patterns for both computation and communication. Although this research focuses on parallelism within FPGAs the primitive operations are examples of fundamentally parallel operations, exposing high levels of optimisation potential through SIMD data-paths. The authors look to extract performance from their FPGA implementation of potentially parallelisable algorithms; their key focus is on the Parzen window technique of Gaussian PDF estimation, as well as K-Means clustering. From their research they develop a concept of pattern-based algorithm decomposition; as a means of exposing potential parallelism within a codebase to individuals with little or no experience with highly parallel code. The authors describe a limitation in the existing development framework, where FPGA based systems are developed on by only highly specialised and skilled developers. This concept is explored by the authors of [11] in which they describe the limitation upon many programmers is the lack of a breadth of libraries which exploit the benefits of GPGPU programming. Within this paper they present the problem of fragmentation within the GPU programming space, with competing standards and APIs resulting in specialists rolling their own solutions to many problems, resulting in minimal code re-use. The current solution to the code reuse problem is the introduction of parallel algorithms within the C++ STL as well as CUDA based libraries such as CuBLAS and Thrust, which aim to provide programmers with a foundation from which to build larger programs, without being forced into having a full understanding of programming on a GPU. Given the power and performance characteristics of GPUs there is a great incentive to move computationally expensive algorithms onto these devices, currently however there are many roadblocks preventing individuals and organisations from exploiting the potential improvements within their codebase, the high barrier to entry can be difficult to overcome, especially with the limited capabilities of identifying how beneficial any redevelopment may actually be.

11

4.4 Static Analysis

Static Analysis is the analysis of the original source statements of a program, it provides a method by which the semantics of a program may be identified. Typically, Static Analysis is utilised in order to identify bugs within a codebase [12], there is a litany of literature which describes a variety of processes through which one can employ Static Analysis to search out and find bugs [13] [14]. These tools rely on parsing the source of a program and identifying the anti-patterns within, alerting developers to their potential mistakes. Additionally, many static analysers provide complexity statistics about the analysed source; it is often the intention of static analysers to provide the developer with a summary of information about their codebase which facilitates further investigation and development. Although there has been a large amount of research into static analysis for the purpose of detecting and eliminating bugs, there is a lack of research into the topic of utilising static analysis for the purpose of performance improvements. Static analysis as a tool is rarely used to facilitate optimisation improvements within a codebase; programmers often utilise profiling in order to determine where to invest development time focused on optimisation. Clearly this presents a divide, where programmers often utilise Static Analysis and Profiling in a competing demands structure for developer time. Within the NVidia High-Productivity CUDA Programming presentation [15] they recommend the utilisation of a process called APOD Assess-Parallelise-Optimise-Deploy. Within the Assessment and Optimisation phase of the process they recommend utilising profilers in order to make determinations about what aspect of the code requires further development. The weakness of profiling however is that it is often very time consuming, additionally it requires an individual developer have an understanding of how to potentially translate sequential serial algorithms into their massively parallel equivalents.

4.5 Summary

TODO:

5   Parallel Computation

Parallel computing is a type of computation whereby many operations are performed simultaneously, as opposed to serial computing where only one operation occurs at a time [16]. Parallel computing has been used as a high performance computing technique for some time, with recent physical limitations on serial processors forcing further development in the parallel world. Modern parallel computing focuses primarily on extracting maximum performance from data level parallelism, the process by which independent processors act on a distributed load of data, often performing SIMD (Signle Instruction Multiple Data) operations.

5.1   SIMD

SIMD describes a computation structure by which many processing units execute the same instruction on multiple data in parallel. SIMD allows for significant computation speed improvements over traditional serial data processing by operating on multiple data at once. The theoretical speed improvements a SIMD processor has over a traditional serial processor can be described by: $speedup \propto min\left(N_{processing\_units}, N_{data}\right)$. In many computing environments N_{data} is significantly larger than $N_{processing\_units}$ simplifying the relationship to merely $speedup \propto N_{processing\_units}$. Thus it is clear that with increasing number of processing units the commputation speed increases. As a result of this relationship devices have been developed which can exploit this fact, most modern CPU's include some form of Vectorised SIMD instruction set. Advanced Vector Extensions (AVX) are an extension to the x86 instruction set which are supported by both AMD and Intel, which utilises SIMD to improve processor performance for highly parallel workloads. GPU's however are far more suited to the task of performing SIMD operations as they often have orders of magnitude more processing units than comporable CPU's.

5.2   Parallel Limitations

One of the largest limitations that concerns parallel computing is data dependency. A data dependency is a situation where operations in the algorithm require data from earlier in the algorithm in order to continue processing. An example situation

would be:

```
1  double mean = 0;
2
3  // Calculate Mean
4  for (size_t i = 0; i < ELEMS; ++i)
5      mean += vec[i]/ELEMS;
6
7  // Data Dependency: Relies on Calculated Mean
8  for (size_t i = 0; i < ELEMS; ++i)
9      vec[i] = abs(vec[i] - mean);
10
11 // Data Dependency: Relies on other value of vec
12 for (size_t i = 0; i < ELEMS; ++i)
13     vec[i] = vec[vec[i]];
```

in this case we have two examples of data dependency, in the first case we are trying to normalise the vector by subtracting the mean. This is an example of a data dependency where a SIMD operation relies on prior information, in this case this will not impact our ability to perform the operations simultaneously, as at no point does the input information rely on potential changes to the output information as a side-effect of this computation. However in the second case, where we re-assign the vector values, there is a data dependency that prevents a Parallel Implimentation from being naivly implemented. `vec[vec[i]]` has a dependency on prior operations performed to `vec[]` which prevents us from processing all elements of this loop simultaneously. There are however classes of problems which may be simply parallelised, these are known as Parallel Primitives.

5.3   Parallel Primitives

Parallel Primitives, or Vector Primitives are operations over a collection of values, there are three operations which constitute these primitives:

- Map

- Reduce

- Scan (Prefix Networks)

These operations all rely on the ability to reformat the problem specification to utilise a computation graph for simultanous calculation.

5.3.1 Map

```
for (size_t i = 0; i < SIZE; ++i)
  out_vec[i] = in_vec[i] * 2;
```

Map operations are the simplest of the three Parallel Primitives, they are simply operations describing a one to one mapping from some input value to some output value, Mapped over a set of values. Map operations have some restrictions upon them about what is considered valid inputs. Operators must have an arity of 1 and the operator must be stateless. If these rules are held then the system will be by definition an LTI system, allowing for a trivial SIMD implementation of the resulting transformation. If however the input arguments do not satisfy these requirements, then the resulting operation will not be well formed, and in most implementations the result will be ill-defined. Map operations have a work complexity of $O(N)$ for CPU implementations and $O(N/k)$ for parallel implementations where k is the number of computational cores available.

5.3.2 Reduce

```
for (size_t i = 0; i < SIZE; ++i)
  k += in_vec[i];
```

Reductions are the next simplest of the Parallel Primitives, they are a mapping from many input values to one output value. Input argument restrictions on Reductions are more strict than on Map operations. Reductions require that the operator and data form a monoid; the Operator must be a binary associative operator, and the set of input values must be closed under that operator. A simple example would be addition, over the Reals. Like with the Map primitive, if the input restrictions are not met, then the operation will not be well formed. Reduce operations have a work complexity of $O(N)$ for CPU implementations and $O(N)$ for parallel implemetations, where the Step Complexity is $O(logN)$. Reductions form a work-efficient operation.

Additional optimisation can occur if the operator is not only associative but also commutative, whereby the gathering of values may occur out of typical oder providing potentially better constant factors.

Figure 1: Serial vs Parallel Sum Reduction Tree

5.3.3 Scan

```
1  out_vec[0] = in_vec[0];
2  for (size_t i = 1; i < SIZE; ++i)
3    out_vec[i] = in_vec[i] + out_vec[i-1];
```

Scans or Prefix Networks are the most complext of the Parallel Primitives. Scans
are a mapping from many input values to many output values. They are a direct
generelisation of a Reduction, where the cumulative intermediate values are main-
tained. Scan operations require the same restrictions for Reductions hold. Scans
are not trivially parallelisable, as there is a data dependency on prior calculated
values, however there are algorithms for performing parallel Scan operations whilst
still retaining work efficiency (a work complexity of $O(N)$) [17].

5.3.4 Linear Algebra

TODO: THIS (NOT A PRIMITIVE, BLAS HOWEVER IS VERY PARALLELIS-ABLE)

5.4 GPU Parallelism

GPUs typically consist of thousands of processing cores, this is significantly greater than the common 4-8 cores found on modern CPUs. GPU architecture relies on numerous simple processing units which engage in SIMD, leveraging the relationship between computational cores and work efficiency.

Figure 2: CPU vs GPU Architecture [18]

Efficient GPU programming like efficient CPU programming requires specialised knowledge, both of the target hardware, and of the paradigm. GPU programming is able to exploit the massively parallel compute archictecture on these devices. Coupling fast global memory, high performance shared memory and numerous local registers GPUs provide all the requirements for exploiting SIMD in a massively parallel space.

6 Clang Integration

CAPA utilises the Clang Libtooling library in order to perform semantic static analysis of any C codebase, via a hook into the clang frontend action.

6.1 Clang Compilation

Clang is the C frontend of the llvm project, and as such it compiles source code into llvm intermediate representation, rather than directly to ASM. This allows Clang to target a common language leveraging the llvm compiler to perform the final compilation to a native binary. A core tenant of the Clang toolset is that Clang is more than just a Compiler, it also a library, which allows for third party individuals to use and extend the Clang framework in a variety of ways. [19] The Clang frontend compilation phase is a 4 step process, where Source code is lexed to tokens, parsed into an AST for programatic manipulation. Semantic Analysis is performed on the AST to identify standards compilance and enforce some optimisations, before being process by the Code Generator which exports LLVM IR. Each of these phases in the frontend action is exposed via a public library which can be utilised by third party tools.

6.2 Benefits of Clang

Clang provides a number of benefits over rolling your own source code analyser. Primarily Clang is one of the most used C language compilers in the world, it is well maintained, well documented, standards compliant and provides first class library support at most levels of compilation. Using existing tooling for the heavy lifting of the project allowed for far more time to be invested in further developing the project, rather than scaffolding the basic fundamentals. Using the Clang libraries also allows for expansion in the future to parsing more than just C files, there is potential to parse any file compilable by Clang, from C++, Obj-C and from Clang4.0 onwards Cuda-C [19].

6.3 Integrating CAPA

The Clang compiler exposes a public library for interfacing with their intermediate compilation stage representations of the original source. For the purpose of this project it was decided to use the exposed AST interface in order to perform the static analysis. Clang provides a number of methods for working with the AST, namely the Visitor and Matcher interfaces. The Visitor library utilises the visitor pattern, and a callback is undertaken upon visiting any node which meets the requirements set forth in the visitor module. This is a useful tool, however it is not as powerful as the Matcher interface, which allows complex grammars to be generated for highly specific, tailored matches. CAPA utilises the ASTMatcher callback interface in order to provide complex generic and extensible traversals of the AST.
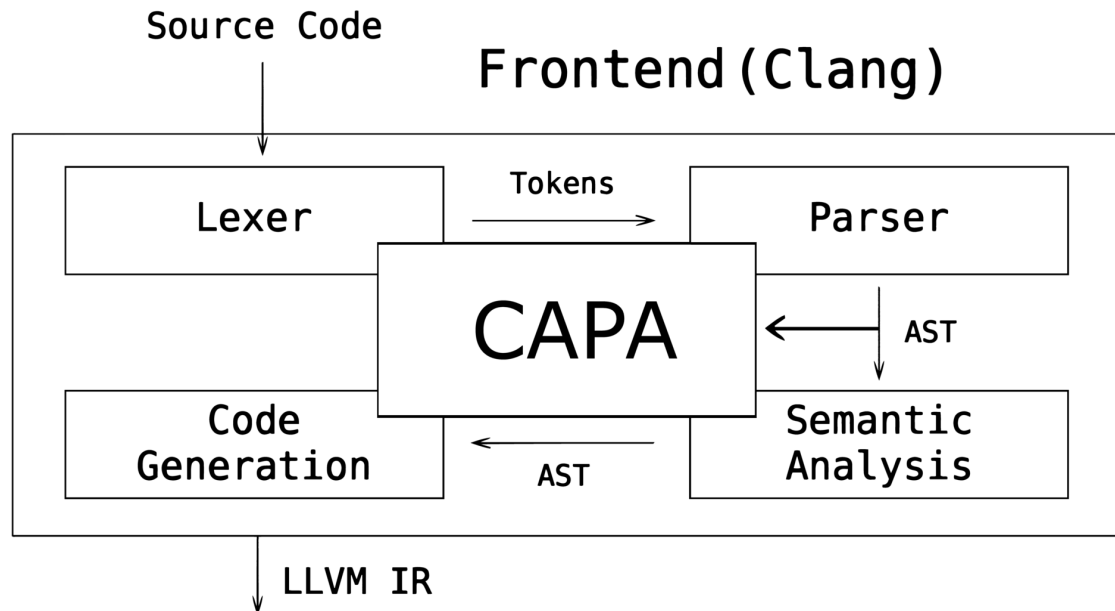
Figure 3: CAPA Hook-In

As CAPA is a fork of OCLint [2] the scaffolding around the Clang integration was already provided. CAPA accepts compile flags via the command line which are passed through to the existing Clang compilation libraries which engage in the frontend action. CAPA hooks into the compilation phase and provides requests after the parser has constructed the AST. At this point the matcher interface begins

to operate on the internal AST, wherein CAPA begins to analyse the provided source files.

7 Project Components

This section aims to provide an in depth look at what has been achieved over the life of the project.

7.1 Static Analyser

7.1.1 Clang Integration

CAPA utilises the Clang Libtooling library in order to perform semantic static analysis of any C codebase. The Clang compiler exposes a public library for interfacing with their intermediate compilation stage representations of the original source. For the purpose of this project as described in section 6.3 it was decided to use the exposed AST interface in order to perform the static analysis. The matcher interface provides a declarative API by which the program searches for regions which satisfy known static requirements. The interface itself is rather unwieldly to use directly.

```
1  auto MapMatcher =
2  forStmt(
3      hasLoopInit(anyOf(
4          declStmt(hasSingleDecl(varDecl(hasInitializer(
5              integerLiteral(anything())))).bind("InitVar"))),
6          binaryOperator(
7              hasOperatorName("="),
8              hasLHS(declRefExpr(to(varDecl(hasType(
9                  isInteger()))).bind("InitVar"))))))),
10     hasCondition(binaryOperator(hasRHS(hasDescendant(
11         declRefExpr().bind("var"))))),
12     hasIncrement(unaryOperator(
13         hasOperatorName("++"),
14         hasUnaryOperand(declRefExpr(to(varDecl(hasType(isInteger()))
15             .bind("IncVar")))))),
16     hasBody(hasDescendant(binaryOperator(
17         hasOperatorName("="),
18         hasLHS(arraySubscriptExpr(
19             hasBase(implicitCastExpr(hasSourceExpression(declRefExpr(to(
20                 varDecl().bind("OutBase")))))),
21             hasIndex(hasDescendant(declRefExpr(to(varDecl(hasType(
22                 isInteger()))).bind("OutIndex"))))))),
23         hasRHS(hasDescendant(arraySubscriptExpr(hasBase(implicitCastExpr(
24             hasSourceExpression(declRefExpr(to(varDecl().bind("InBase")))))),
25                 hasIndex(hasDescendant(declRefExpr(to(varDecl(hasType(
26                     isInteger()))).bind("InIndex")))))))),
27         unless(hasDescendant(arraySubscriptExpr(hasDescendant(binaryOperator())))
             ))
28         .bind("Assign")))).bind("Map");
29
30
```

```
31  addMatcher(MapMatcher);
```

As a result I created a matcher combinator library for simplification purposes.

7.1.2   ASTMatcher Combinator Library

In order to simplify the matchers and prevent them from becoming unmanagebly
large, I designed a lambda based combinator library for creating complex AST-
Matchers. Utilising C++14 auto lambda return type declarations I was able to
construct a number of higher order combinators which can be combined together
to construct more complex matchers with less ambiguity. A simple example:

```cpp
1   namespace CAPA {
2
3   // Variable Binding Combinators
4   auto VarBind = [](std::string binding)
5   {
6       return declRefExpr(to(varDecl().bind(binding)));
7   };
8
9   auto DVarBind = [](std::string binding)
10  {
11      return hasDescendant(VarBind(binding));
12  };
13
14  auto VectorBind = [](std::string binding)
15  {
16      return arraySubscriptExpr(
17              hasBase(DVarBind(binding + "Base")),
18              hasIndex(DVarBind(binding + "Index")));
19  };
20
21  auto MatrixBind = [](std::string binding)
22  {
23      return arraySubscriptExpr(
24              hasBase(hasDescendant(arraySubscriptExpr(
25                  hasBase(DVarBind(binding + "Base")),
26                  hasIndex(DVarBind(binding + "Row"))))),
27              hasIndex(DVarBind(binding + "Column")));
28  };
29
30  // Loop Binding Combinators
31  auto LoopInit = [](std::string level)
32  {
33      return anyOf(
34          declStmt(hasSingleDecl(varDecl(hasInitializer(
35              integerLiteral(anything()))).bind("InitVar" + level))),
36          binaryOperator(
37              hasOperatorName("="),
38              hasLHS(VarBind("InitVar" + level))));
39  };
40
41  auto LoopIncrement = [](std::string level)
42  {
```

```
43    return anyOf(
44        unaryOperator(anyOf(
45            hasOperatorName("++"),
46            hasOperatorName("--")),
47          hasUnaryOperand(VarBind("IncVar" + level))),
48        binaryOperator(anyOf(
49            hasOperatorName("+="),
50            hasOperatorName("-=")),
51          hasLHS(VarBind("IncVar" + level)),
52          hasRHS(expr().bind("Stride" + level))));
53  };
54
55  auto ForLoop = [](std::string binding, std::string level, auto injectBody)
56  {
57    return forStmt(
58        hasLoopInit(LoopInit(level)),
59        hasCondition(binaryOperator(hasRHS(expr().bind(binding + "CondRHS"))
                )),
60        hasIncrement(LoopIncrement(level)),
61        hasBody(injectBody)).bind(binding);
62  };
63
64  } // End Namespace CAPA
```

This combinator library again is easily extensible and as further needs arose I increased its complexity and breadth. Ultimately it provides a simpler way of interfacing with the Clang AST Matcher library through safer higher order functions. This combinator library would not be possible with earlier versions of C++, as it has a reliance on auto return types to prevent template instantiation stack errors. The AST Matcher interface heavily relies on template meta-programming in order to ensure a clean typesafe interface, the combinator library I designed extends that, yet still maintains complete statically verifiable interfaces that are type correct.

```
1  auto left = VectorBind("Out");
2  auto right = hasDescendant(VectorBind("In"));
3  auto unless = hasDescendant(arraySubscriptExpr(hasDescendant(binaryOperator())
    ));
4
5  auto body = hasDescendant(BinaryOperatorBindUnless("=", "Assign", left, right,
    unless));
6
7  auto MapMatcher = ForLoop("Map", "", body);
8
9  addMatcher(MapMatcher);
```

This version is clearly far simpler to understand and to modify, whilst still achieving the same callback results as the original version. This demonstrates the power of the Matcher Combinator interface, as well as the ease of use.

7.1.3 Pattern Recognition

7.1.3.1 Matching By utilising the OCLint tool the scaffoling around the libtooling had already been provided, simplifying the interface between pattern matching and reporting. There was still a significant rewrite of most of the interface, however the heirachy and design philosophy was clear. The actual matching of potentially parallel sections of code relies on a few assumptions about the nautre of algorithms which may be efficiently implemented on a GPU.

- Little prior data dependency

- A large number of elements require processing

- Minimal control flow is required

Given these assumptions, in combination with the knowledge of problem spaces that the GPU is highly performant in:

- Map Operations

- Reduce Operations

- Scan (Prefix Network) Operations

- Matrix Operations

it became clear that identifying components of the codebase that exhibited similarities to these cases would be important.

7.1.3.2 Callback Even though the matcher can be highly specific, there is still the requirement of the callback. The callback is responsible for identifying whether the matched pattern is actually representative of the expected pattern, or whether it is infact a potential false positive. Additionally the callback is responsible for logging the detected pattern as well as relevant information, such as the number of elements being manipulated, for use by the reporter module. This is explained in more detail in section 8.

7.1.4 Benchmark Integration

A key aspect of CAPA is the integration of the benchmarking component with the detected outcomes to provide more detailed performance statistics during the reporting phase.

7.1.4.1 JSON Parsing CAPA relies on performance benchmarks being generated by the benchmarking tool, this reports back information in a JSON form which is read and parsed by CAPA into a useable form. The JSON for Modern C++ library was used [20]. JSON was chosen as it is a human readable format which is supported by a large number of tools, additionally the benchmarking libary utilised provided a JSON exporter, simplifying the integration of the two tools.

7.1.4.2 Internal Representation Internally the benchmarking information relies on this construct:

```cpp
class BenchmarkSet
{
public:
    BenchmarkSet(std::string benchmarkLocation);
    BenchmarkSet();
    bool Exists(std::string operation) const;
    double Speedup(std::string operation, std::size_t dimension) const;
    double Speedup(std::string operation) const;
    std::tuple<double, double> GetResult(std::string operation, std::size_t
        dimension) const;
    std::tuple<double, double> LowerUpper(std::string operation, std::size_t
        dimension) const;

private:
    std::map<std::string, std::map<std::size_t, std::tuple<double, double>>>
        benchmarks;
    static const std::map<const std::string, const std::size_t> VectorFixtures;
    static const std::map<const std::string, const std::size_t> MatrixFixtures;
};
```

which simply provides a simple concise manner in which to interface with STL containers for the purpose of passing around the parsed benchmarking information. The BenchmarkSet object is responsible for handling all requests for benchmark information, including providing theoreticals if no benchmark JSON file exists. This information is then used by the reporter module to provide the end user with further information about potential optimisations within their analysed codebase.

7.1.5 Reporting

The reporter module is dynamically loaded at runtime to provide the end user with readily swappable components, this is a continuation of the design decision implemented by the designers of OCLint. The reporter module is reponsible for accepting all recognised patterns and providing an output. Reporters are also responsible for utilising performance criteria calculated by the benchmark suite and integrating this information into the output report. CAPA currently has only one reporter, the text reporter which provides coloured output for ANSI terminals.

7.2 Benchmarks

A key component of this project is the ability of CAPA to provide performance characteristics for a given system, thus a benchmarking suite was developed to gauge comparitive CPU and GPU performance for known problem sets.

7.2.1 Benchmark Structure

The benchmarking suite is designed using the hayai[21] library for benchmark automation, the hayai library provides a macro interface similar to googleTest, this was used to benchmark both Host and Device performance.

7.2.2 Libraries Benchmarked

Three libraries were benchmarked through this process, in order to emulate industry practice. For host and device side parallel primitives the CUDA Thrust[4] library was used. For Matrix Multiplication the Eigen[3] library was used for the host side implementation, and CuBLAS[5] was used for device side implementation.

7.2.3 Preperation

Template metaprogramming was utilised to ensure a generic and extensible testing framework was generated. A monolithic benchmark object is responsible for setting up and running the individual operations, which are accessed through a benchmarking fixture via object composition.

7.2.4 Implementations

All benchmarks were run with floating point precision, however all benchmarks are templated over their implementation type, so benchmarks are lower or higher precisions are capable.

Additionally all device benchmarks include onload and offload time for memory management, and no device benchmarks utilise explicit streaming or pinned memory.

The benchmark code was compiled at with no optimisations and with `-fno_vectorize` enabled, to ensure all host code was purely sequential.

7.2.4.1 Map The map benchmark simply implemented a negation across a number of random elements.

Host The host map operation is directly implemented by thrust:

```
1    void host_map()
2    {
3        thrust::transform(thrust::host, H.begin(), H.end(), H.begin(), thrust::
             negate<T>());
4    }
```

Device The device map operation is directly implemented by thrust, however the Device Vector is first intisialised by the Host vector before being operated on.

```
1    void device_map_onload()
2    {
3        D = H;
4        thrust::transform(thrust::device, D.begin(), D.end(), D.begin(), thrust
             ::negate<T>());
5    }
```

7.2.4.2 Reduce The reduce bencmark is parameterised over a binary operation with a known identity, this is enforced by the use of template metaprogramming and a simulated closure object via a C++ Functor. The reduction benchmark traverses the vector searching for the minimum value.

Host The host operation relies on thrust for the computation.

```
1    template <template <class> class Functor>
2    T host_reduce(Functor<T> &binaryOp)
3    {
4        return thrust::reduce(thrust::host, H.begin(), H.end(), binaryOp.
             identity, binaryOp);
5    }
```

Device   The device operation relies on thrust for the computation.

```
1    template <template <class> class Functor>
2    T device_reduce_onload(Functor<T> &binaryOp)
3    {
4        D = H;
5        return thrust::reduce(thrust::device, D.begin(), D.end(), binaryOp.
             identity, binaryOp);
6    }
```

7.2.4.3   Scan (Prefix Network)   The scan benchmark is also parameterised over a binary operation with a known identity, again this enforced by the C++ template engine. The scan benchmark is also traversing the vector searching for the minimum value, storing all intermediate results.

Host   The host operation relies on thrust for the computation.

```
1    template <template <class> class Functor>
2    void host_scan(Functor<T> &binaryOp)
3    {
4        thrust::inclusive_scan(thrust::host, H.begin(), H.end(), H.begin(),
             binaryOp);
5    }
```

Device   The device operation relies on thrust for the computation.

```
1    template <template <class> class Functor>
2    void device_scan_onload(Functor<T> &binaryOp)
3    {
4        D = H;
5        thrust::inclusive_scan(thrust::device, D.begin(), D.end(), D.begin(),
             binaryOp);
6    }
```

7.2.4.4   Dense Matrix Multiplication   Dense matrix multiplication benchmarks were implemented via Eigen for the host side, and CuBLAS for the device

side. Thrust was used for memory management on the device side.

Host   The host operation relies on Eigen for the computation. Eigen provides overloaded operators for matrix operations.

```
1   void host_matrix_mult()
2   {
3       eMatrixC = eMatrixA * eMatrixB;
4   }
```

Device   The device operation relies on Thrust for memory management and CuBLAS for the actual matrix multiplication.

```
1    void cuBLAS_prepare()
2    {
3        matrixA.resize(hA * wA);
4        matrixB.resize(wA * wB);
5        matrixC.resize(hA * wB);
6        cublasCreate(&handle);
7        device_matrix_mult();
8        device_matrix_mult();
9    }
10
11   void device_matrix_mult()
12   {
13       thrust::device_vector<T> d_A = matrixA;
14       thrust::device_vector<T> d_B = matrixB;
15       thrust::device_vector<T> d_C(hA * wB);
16
17       cuBLAS_mmul(thrust::raw_pointer_cast(&d_A[0]),
18                   thrust::raw_pointer_cast(&d_B[0]),
19                   thrust::raw_pointer_cast(&d_C[0]),
20                   hA, wA, wB);
21
22       matrixC = d_C;
23   }
24
25   void cuBLAS_mmul(const float *matA, const float *matB, float *matC, const
         size_t m, const size_t k, const size_t n)
26   {
27       size_t lda   = m;
28       size_t ldb   = k;
29       size_t ldc   = m;
30       const float alf   = 1;
31       const float bet   = 0;
32       const float *alpha = &alf;
33       const float *beta = &bet;
34
35       // Do the actual multiplicatio
36       cublasSgemm(handle, CUBLAS_OP_N, CUBLAS_OP_N, m, n, k, alpha, matA, lda,
             matB, ldb, beta, matC, ldc);
37   }
38
```

29

```
39    void cuBLAS_mmul(const double *matA, const double *matB, double *matC, const
          size_t m, const size_t k, const size_t n)
40    {
41        size_t lda   = m;
42        size_t ldb   = k;
43        size_t ldc   = m;
44        const double alf  = 1;
45        const double bet  = 0;
46        const double *alpha = &alf;
47        const double *beta = &bet;
48
49        // Do the actual multiplicatio
50        cublasDgemm(handle, CUBLAS_OP_N, CUBLAS_OP_N, m, n, k, alpha, matA, lda,
              matB, ldb, beta, matC, ldc);
51    }
```

7.2.5   Typesafe CUDA Primitives

Additionally in order to increase code re-use I've attempted to implement the
concept of Type-Classes into CUDA C++. This allows for the single case of
higher order functions such as Map, Reduce and Scan. This implementation utilises
template meta-programming and templated type alias's to produce type safe higher
order polymorphism within CUDA compute kernels.

```
1   template <typename T>
2   using uCat = T(*)(T);
3
4   template <typename T>
5   using mCat = T(*)(T, T);
6
7   __device__ int square(int a)
8   {
9       return a * a;
10  }
11
12  __device__ int mult(int a, int b)
13  {
14      return a * b;
15  }
16
17  template <typename T>
18  __device__ T mult(T a, T b)
19  {
20      return a * b;
21  }
22
23  template <typename T, mCat<T> func>
24  __global__ void biMapKernel(T *a, T *b, T *c, size_t size)
25  {
26      size_t i = threadIdx.x + blockIdx.x * blockDim.x;
27      if (i < size)
28          c[i] = func(a[i], b[i]);
29  }
30
```

```
31  template <typename T, uCat<T> func>
32  __global__ void MapKernel(T *a, T *c, size_t size)
33  {
34      size_t i = threadIdx.x + blockIdx.x * blockDim.x;
35      if (i < size)
36          c[i] = func(a[i]);
37  }
```

This code segment demonstrates typeclass instances for operations over both *Functors* and *BiFunctors*, the *Functor* typeclass is the alias `uCat`, a function which accepts a single input of type `T` and returns a single output of type `T`. The second typeclass definition is the `mCat` defition, defining the monoidal typeclass requirement of *mConCat*. This is any function that accepts 2 inputs of type `T` and returns an output of type `T`. These typeclasses are then used to validate the parametric polymorphism of the `biMapKernel`, which implements a polymorphic bimap operation, the key aspect of the *bifunctor* typeclass, as well as the common map kernel, which is the single requirement of the *functor* typeclass. This parametric polymorphism provides an easy to use, clear and concise framework from which to build larger GPU benchmarking routines, by utilising the higher order nature of the GPU primitives. To summarise this essentially provides a type safe way to write less code.

8   Case Study

In order to fully understand CAPA it's important to explore a case study demonstrating the process of analysing code under test.

```c
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <math.h>

///CAPA:IGNORE
void random_fill(float *starting_vec, size_t size){
    for (size_t i = 0; i < size; ++i)
        starting_vec[i] = (float) rand();
}

/** This Function is reponsible for reshaping a vector
 * into a matrix.
 * CAPA:IGNORE
 */
void reshape2mat(float *in_vec, float *out_vec[], size_t dim){
    for (size_t i = 0; i < dim; ++i)
        for (size_t j = 0; j < dim; ++j)
            out_vec[i][j] = in_vec[i*dim + j];
}

void reshape2vec(float *out_vec, float *in_mat[], size_t dim){
    for (size_t i = 0; i < dim; ++i)
        for (size_t j = 0; j < dim; ++j)
            out_vec[i*dim + j] = in_mat[i][j];
}

/// CAPA:IGNORE
void mmult(float **A, float **B, float **C, size_t dim);

int main()
{
    // Toy Example for Test example Case.
    const size_t ELEMS = 1000*1000;
    float starting_vec[ELEMS];
    time_t t;

    srand((unsigned) time(&t));
    random_fill(starting_vec, ELEMS);

    // Divide all points by 2 and add 4 for later stage
    for (size_t i = 0; i < ELEMS; ++i){
        starting_vec[i] /= 2;
        starting_vec[i] += 4;
    }

    // Calculate mean
    float k = 0;
    for (size_t i = 0; i < ELEMS; ++i)
        k += starting_vec[i]/ELEMS;

    // Renormalise all values and compute cumulative sum
    float cum_sum[ELEMS];
```

```
54      cum_sum[0] = starting_vec[0]/k;
55      for (size_t i = 1; i < ELEMS; i++)
56          cum_sum[i] = starting_vec[i]/ELEMS + cum_sum[i-1];
57
58      // Prepare for and perform Matrix Mult
59      const size_t dim = sqrt(ELEMS);
60      float cum_sum_mat[dim][dim];
61
62      reshape2mat(cum_sum, (float **) cum_sum_mat, dim);
63      mmult((float **) cum_sum_mat, (float **) cum_sum_mat, (float **) cum_sum_mat
            , dim);
64      reshape2vec(cum_sum, (float **) cum_sum_mat, dim);
65
66      // reduce only even values, but do it with conditional
67      // and deliberately prevent it being caught
68      k = 0;
69      for (size_t i = 0; i < ELEMS; ++i)
70          if (!(i % 2))
71              k = k + cum_sum[i + 1 - 1];
72
73      return k;
74  }
75
76  void mmult(float **A, float **B, float **C, size_t dim){
77      for (size_t i = 0; i < dim; ++i)
78          for (size_t j = 0; j < dim; ++i)
79              for (size_t k = 0; k < dim; ++i)
80                  C[i][j] += A[i][k] * B[k][j];
81  }
```

```
CAPA Report

Summary: TotalFiles=1 Files With Improvements=1

/home/james/Projects/LatexDocs/DesignDoc/Code/CodeUnderTest.cpp:22:1
Pattern: Vectorisable Function Declaration Priority: 3 Info: Function Declaration
void reshape2vec(float * out_vec, float ** in_mat, size_t dim);

/home/james/Projects/LatexDocs/DesignDoc/Code/CodeUnderTest.cpp:76:1
Pattern: Vectorisable Function Declaration Priority: 3 Info: Function Declaration
void mmult(float ** A, float ** B, float ** C, size_t dim);

/home/james/Projects/LatexDocs/DesignDoc/Code/CodeUnderTest.cpp:42:5
Pattern: Map Priority: 3 Info: Stride Size: 1. Number of Elements: 1000000.
Potential Speedup: 158.03 ~ 140.51
for (size_t i = 0; i < ELEMS; ++i){
        starting_vec[i] /= 2;
        starting_vec[i] += 4;
    }

/home/james/Projects/LatexDocs/DesignDoc/Code/CodeUnderTest.cpp:49:5
Pattern: Reduce Priority: 2 Info: Stride Size: 1. Number of Elements: 1000000.
Potential Speedup: 30.30 ~ 34.33
for (size_t i = 0; i < ELEMS; ++i)
        k += starting_vec[i]/ELEMS

/home/james/Projects/LatexDocs/DesignDoc/Code/CodeUnderTest.cpp:55:5
Pattern: Scan Priority: 2 Info: Stride Size: 1. Number of Elements: 1000000.
Potential Speedup: 19.72 ~ 30.11
for (size_t i = 1; i < ELEMS; i++)
        cum_sum[i] = starting_vec[i]/ELEMS + cum_sum[i-1]

/home/james/Projects/LatexDocs/DesignDoc/Code/CodeUnderTest.cpp:77:5
Pattern: Matrix Multiplication Priority: 1 Info: A Matrix Multiply
Potential Speedup: 16.72 ~ 229.00
for (size_t i = 0; i < dim; ++i)
        for (size_t j = 0; j < dim; ++i)
            for (size_t k = 0; k < dim; ++i)
                C[i][j] += A[i][k] * B[k][j]

/home/james/Projects/LatexDocs/DesignDoc/Code/CodeUnderTest.cpp:23:5
Pattern: Vectorisable region Priority: 5 Info: Generally vectorisable region of code
for (size_t i = 0; i < dim; ++i)
        for (size_t j = 0; j < dim; ++j)
            out_vec[i*dim + j] = in_mat[i][j]


[CAPA v0.10.2]
```

Figure 4: CAPA Generated Report

Please note that the cast study requires knowledge of the Clang AST, the output
AST dump for the file under test can be found in the Appendix at section 11.1.

8.1 Setup

The initial setup is mostly taken care of by existing Clang and LLVM libraries. Clang Libtooling is responsible for parsing and lexing the source file, CAPA merely passes the arguments through to the correct Clang libraries. During the setup CAPA dynamically loads Rules for AST analysis and Reporters for reporting. Once the rules and reporters have been loaded, CAPA loads the benchmark information for use in the reporting phase.

Loading of rules consists of two phases. Rule Generation and Collection. Rule Generation is the setup process whereby the AST Matchers are created, and the resulting matcher is then forwarded onto Clang with a callback function for later processing. The rules are then collected by CAPA for dispatching callbacks correctly.

Once the file has been loaded by the Clang front-end, parsed, and the AST constructed, the Libtooling library begins to traverse the AST searching for a match.

8.2 Matching

When Clang finds a portion of the AST which matches the requirements described by the rules, the callback function is envoked and the relevant rule must process the matching AST Node.

For our code under test, the first callback that occurs is not shown in the report, this is due to the tag which tells CAPA to ignore that function. Further information can be found here 8.3.

8.2.1 Match 1: Vectorisable Function Declaration

The first match to be reported is the result of the function declaration on line 22. The pattern identified is a `Vectorisable Function Declaration`. This is the simplest of all the parallel matchers. A function declaration is deemed to be potentially vectorisable based purely on the type information available. Vectorisable functions require an arity of at least two, with one argument being a pointer or array type, and the other argument being a `size_t`. This is the most speculative of the patterns identified by CAPA, as there is very little information to work with in a function

declaration, however functions that operate on vectors in C require both a pointer to the vector, and some information about the size of the vector, thus with that limited information we can construct a matcher.

The matcher for this rule is described by:

```
auto Matcher = functionDecl(allOf(
                anyOf(hasAnyParameter(hasType(arrayType())),
                    hasAnyParameter(hasType(pointerType()))),
                hasAnyParameter(hasType(asString("size_t"))))).bind("Function"
                    );
```

The grammar here is quite simple, the matcher is requesting a callback if a function declaration is found which has a parameter that is either an array type or a pointer type, and has any parameter which is a size_t. If this is found it is to be bound by the name Function and the callback will provide the bound nodes and additional AST context.

8.2.1.1 Match 2: Vectorisable Function Declaration

The second match is also a Vectorisable Function Declaration. This match however is a the result of catching a function declaration on line 76. The process by which this function is found is no different from the prior example. Note however that the types are different, and in this case the matcher is in fact catching types with two levels of indirection. This rule is general enough that it is capable of catching arbitrary levels of indirection.

8.2.2 Match 3: Map Operation

The third match to be reported is a map operation, which is identified on line 42. The reported information for Map operations is more thorough than the information provided for Vectorisable Function Declaration, this is because more information is available to the callback due to a stronger grammar.

```
// Matches on For Loops with counter initialised in the init, with an array
    element
// assignment within the body of the loop
auto left = VectorBind("Out");
auto right = HasDescendant(VectorBind("In"));
auto unless = HasDescendant(arraySubscriptExpr(HasDescendant(binaryOperator()))
    );

auto body = HasDescendant(BinaryOperatorBindUnless("=", "Assign", left, right,
    unless));
```

```
 8   auto body2 = HasDescendant(BinaryOperatorBindAll("Assign", left,
         NumericLiteralBind("Literal")));
 9
10   auto ForMatcher = FunctionWrap(ForLoop("Map", "", anyOf(body, body2)));
11
12   auto WhileMatcher = FunctionWrap(WhileLoop("Map", "", body));
```

The Map matcher utilises the combinator library designed for this project in order
to simplify the top level expression. Our matcher is looking to identify loops which
have an assignment where there is a vector on both sides of the operator, or loops
which have a compound binary operator and some numeric literal. This describes
most of the semantic restrictions a map operation enforces upon a programmer
however in order to extract performance information extra bindings are required.

The report states that there are 1 million elements to be processed, in order
to extract this information, the `ForLoop` combinator provides a binding on the
conditional which is exposed in the callback. Similarly the number of elemts is
also bound to by the `ForLoop` combinator with the respective node being exposed
to the callback.

Upon a successful match, the callback function is called with the results struc-
ture. The results structure contains all the bound nodes, as well as the AST
Context. These tools are utilised in the callback to filter out false positives, cases
where the grammar of the matcher specification is not strict enough to ensure
only valid cases are matched. Within the callback the results are re-organised into
a simple class which defines a few basic operations. In order for a match to be
validated as truly representative of a Map operation, the results must be verified.
This is done through the `MapInfo.isMap()` method, which confirms that the bound
index's of the input and ouput vectors are related without a data dependency.

Beyond validating matches, the callback is responsible for retrieving and pro-
viding information to the reporters about the known quantities of the region.
That is to say that the callback is responsible for retrieving whatever information
is available from the now exposed bound nodes. In order to calculate the number
of elements to be processed, the right hand side of the loop condition is constant
folded to an integer. If this is possible and there is a result, that information
is then used as the number of elements to process. Similarly the stride size in
the report is 1. This is also bound by the `ForLoop` combinator, which binds to
the loop increment. This binding is then matched against typical loop increment

expressions, and where possible constant folding extracts the stride size.

This information is then passed along to the reporter in order to provide performance metrics.

8.2.3 Match 4: Reduce Operation

The fourth match to be reported is a reduction, which occurs on line 49. Like the Map report, the reduction provides more information about potential improvements. The reduction in the code under test is the calculating of a mean, by summing the contents of the `starting_vec` divided by the number of elements. This is cleary a many to one operation which satisfies the requirements of parallel reductions.

```
auto left = VarBind("Acc");
auto right1 = HasDescendant(VectorBind("In"));
auto right2 = allOf(DVarBind("AccRHS"), HasDescendant(VectorBind("In")));

auto body = anyOf(HasDescendant(BinaryOperatorBindReduceAll("Assign", left,
    right1)),
                  HasDescendant(BinaryOperatorBindReduce("Assign", left, right2))
                  );

auto ForStmtReduceMatcher = FunctionWrap(ForLoop("Reduce", "", body));
auto WhileStmtReduceMatcher = FunctionWrap(WhileLoop("Reduce", "", body));
```

In this case the matcher constructed to capture the reductino is different from the map operation, again however it reads like english. The matcher is looking for a loop which in the body has either an assignment or compound assignment operator, with a single value on the left, and a vector on the right. Once this is matched, Clang passes the relevant information back via the callback which proceeds to validate whether the region of the AST is a viable for parallelisation.

8.2.4 Match 5: Scan Operation

The fifth match reported is a scan operation, which occurs on line 55 of the source file. Again like the map and reduction reports, the scan report also provides information about the number of elements and stride size where available. The matcher is described by:

```
auto left = VectorBind("Out");
auto right = forEachDescendant(VectorBindScan("In"));
auto body = anyOf(HasDescendant(BinaryOperatorBind("=", "Assign", left, right))
    ,
```

```
4                    HasDescendant(BinaryOperatorBindAll("Assign", left, right)));
5
6   auto ScanMatcher = FunctionWrap(ForLoopNoParentLoop("Scan", "", body));
```

8.2.5   Match 6: Matrix Multiplication

The sixth match reported is the matrix multiplication that occurs on line 77. Matrix multiplication is not one of the parallel primitives, however it exhibits massive speed improvements when processed on a GPU compared to a CPU. The matcher for the matrix multiplication is described by:

```
1   auto LoopBody =
2       HasDescendant(
3           binaryOperator(
4               hasOperatorName("+="),
5               hasLHS(MatrixBind("Out")),
6               hasRHS(binaryOperator(
7                   hasOperatorName("*"),
8                   hasLHS(HasDescendant(MatrixBind("Left"))),
9                   hasRHS(HasDescendant(MatrixBind("Right"))))
10      )));
11
12  auto MatrixMultMatcher =
13      FunctionWrap(
14          ForLoop("MatrixMult", "0", HasDescendant(
15              ForLoop("", "1",HasDescendant(
16                  ForLoop("", "2", LoopBody))))));
```

This matcher is slightly different from the previous examples, as it requires the nesting of multiple loops. It still reads like plain english however. The matcher is looking for a for loop which has a child that is a forloop which has a child that is also a forloop that has a compound addition assignment within the body, on the left hand side of the compound assignment there is a Matrix, which is defined as a double indexed vector, and on the right hand side there is a scalar multiplication of two matrix elements. Referring to the AST breakdown in figure 5, we can see directly how the tree is matched.

Figure 5: Matrix Multiplication AST Match Breakdown

8.2.6 Match 7: Vectorisable Region

The seventh match reported is a generally vectorisable region has been detected. This pattern is the most general of those reported by CAPA. The AST is queried for regions that are considered generally vectorisable, which are then rejected if they are already tagged as a more specific operation. The matcher for this pattern is:

```
auto unless = anyOf(
            HasDescendant(ifStmt(anything())),
            HasDescendant(switchStmt(anything())),
            HasDescendant(gotoStmt(anything())),
            hasAncestor(forStmt(anyOf(
                hasAncestor(forStmt(anything())),
                hasAncestor(whileStmt()),
                hasAncestor(doStmt())))),
            hasAncestor(whileStmt(anyOf(
                hasAncestor(forStmt(anything())),
                hasAncestor(whileStmt(anything())),
                hasAncestor(doStmt(anything()))))),
            hasAncestor(doStmt(anyOf(
                hasAncestor(forStmt(anything())),
                hasAncestor(whileStmt()),
                hasAncestor(doStmt())))
            );

auto body = anything();

auto ForMatcher = FunctionWrap(ForLoopUnless("Loop", body, unless));
auto WhileMatcher = FunctionWrap(WhileLoopUnless("Loop", body, unless));
```

This matcher simply looks for a loop of any type which contains no control flow within it. Additionally the matcher requires that any ancestor loops also not contain control flow in order to minimise the risk of a data dependency. The operation on line 23 can not be described by any of our GPU primitives, yet it still meets the requirements for code to be potentially parallelised. Note however in this scenario it would be unwise to parallelise this code, as it is merely memory operations with very minimal computational work per element.

8.3 Tagged Regions

Within the source code there are regions marked `///CAPA:IGNORE` and `/** CAPA: IGNORE */` these alert CAPA that the region is not to be reported on. CAPA achieves this through Clang providing the structured comments interface via declaration statements. The AST Context preserves special comments indicated by

the triple slash or double asterisk, allowing CAPA to hook into the context and retrieve comments assigned to declarations. This is achieved via the following:

```cpp
bool markedIgnore(const Decl *d, SourceManager &sm)
{
    if (!d) { return false; }
    const RawComment *rc = d->getASTContext().getRawCommentForDeclNoCache(d);
    if (rc)
    {
        SourceRange range = rc->getSourceRange();

        PresumedLoc startPos = sm.getPresumedLoc(range.getBegin());
        PresumedLoc endPos = sm.getPresumedLoc(range.getEnd());

        std::string raw = rc->getBriefText(d->getASTContext());
        return (raw.find("CAPA:IGNORE") != std::string::npos);
    }
    return false;
}
```

This is called from within the callback of the matchers to determine whether the region should be processed further, or aborted. The callback is expected to pass the nearest parent declaration to the function in order to determine whether the region should be ignored or not. The current implementation ignores at the function declaration level.

Within the code under test there are three declarations which are marked to ignore. If we extract those segments and remove the ignore directive, CAPA includes them in the report, resulting in this summary:

```
CAPA Report

Summary: TotalFiles=1 Files With Improvements=1

/home/james/Projects/LatexDocs/DesignDoc/Code/CodeUnderTestNoIgnore.cpp:6:1
Pattern: Vectorisable Function Declaration Priority: 3 Info: Function Declaration
void random_fill(float * starting_vec, size_t size);

/home/james/Projects/LatexDocs/DesignDoc/Code/CodeUnderTestNoIgnore.cpp:11:1
Pattern: Vectorisable Function Declaration Priority: 3 Info: Function Declaration
void reshape2mat(float * in_vec, float ** out_vec, size_t dim);

/home/james/Projects/LatexDocs/DesignDoc/Code/CodeUnderTestNoIgnore.cpp:17:1
Pattern: Vectorisable Function Declaration Priority: 3 Info: Function Declaration
void mmult(float ** A, float ** B, float ** C, size_t dim);

/home/james/Projects/LatexDocs/DesignDoc/Code/CodeUnderTestNoIgnore.cpp:7:5
Pattern: Vectorisable region Priority: 5 Info: Generally vectorisable region of code
for (size_t i = 0; i < size; ++i)
        starting_vec[i] = (float) rand()

/home/james/Projects/LatexDocs/DesignDoc/Code/CodeUnderTestNoIgnore.cpp:12:5
Pattern: Vectorisable region Priority: 5 Info: Generally vectorisable region of code
for (size_t i = 0; i < dim; ++i)
        for (size_t j = 0; j < dim; ++j)
            out_vec[i][j] = in_vec[i*dim + j]


[CAPA v0.10.2]
```

Figure 6: Report without ignored regions

which now includes the previously ignored function declarations, and the vectorisable regions within them.

8.4 Report Generation

The report generation is the final part of the CAPA analysis. All patterns have been identified and filtered, the reporter is merely responsible for disaplying the information. Currently there is only a text reporter which produces coloured output for ANSI terminals. The reporter module interfaces with the pattern information provided by the rule set in order to report information that has been gathered by the rules. This information is then used by the reporter to interface with the benchmark information to produce performance characteristics which are displayed if available. In cases where performance metrics are not available, no information is displayed.

9 Re-evaluation of Initial Goals

Considering the project has reached completion, it is important to review the requirements initially set forth in order to ascertain whether the objectives of the project have been met. My final year project involved a large amount of exploratory work, and as the project developed different area's became more important to the final deliverable than were anticipated at the beginning. For the design document submitted in week 12 of semester one, a review of my initial requirements was undertaken, and a re-evaluation of their status within the project was conducted. Since then further developments of the project demanded more time and focus be spent on other aspects, as such it is worth re-evaluating the midway evaluation, and where the project stands at completion. The key area's of the project were.

- GPU Benchmark Development

- Algorithm Analytics Development

- Optimisation Analytics Development

which then allows the following breakdown..

9.1 GPU Benchmark Development

As described earlier, the importance of developing working GPU benchmarking code for known problem classes allows for better analytics and reporting in the serial algorithm analysis portions. This therefore was a key aspect of satisfactorily completing the project. The GPU benchmark development had a number of requirements that describe what the project necessitates.

9.1.1 Requirements

9.1.1.1 [FR.003] The program shall run developed benchmark algorithms to further analytical information.

This requirement relates directly to the overall aim of the project, which is described in the requirements just proceeding this. As the project currently stands there is a completed benchmark suite covering all of the GPU primatives in a highly

extensible and meta-programmable fashion. Each of the following benchmarks has been implemented.

- Peak Map

- Peak Fold/Reduce

- Peak Scan

- Peak Matrix Multiplication

Additionally it is trivial to extend the benchmark suite to cover other aspects of GPU performance.

9.1.1.2 [OA.001] The program shall run custom benchmark algorithms to identify GPU performance.

This requirement was met. In order to satisfy this requirement I produced benchmarking code for GPU performance in problem sets that are known to be performant on a GPU. It was my original intention that as well as benchmarks that are known to be performant, that I would also write benchmarks that may naively appear to be performant, yet further inspection demonstrates that they are not in fact performant. This was a rather large task that did not relate back to the key intentions of the project. It was disappointing that constraints prevented the full exploration of this concept, however as the project developed it was clear that focusing on the positive performance aspects would provide a significantly better final product.

9.1.1.3 [OA.002] The program shall work on all CUDA devices.

After careful consideration this requirement was relaxedi at the midpoint of the project. It was determined to be far too strict. When writing the original requirements analysis I was not as familiar with the CUDA toolchain as I was at the midpoint of the project, and as such it is now apparent that writing Compute Capability agnostic code is a very difficult feat. In order to best satisfy the other requirements of this project I decided to limit benchmarking code to work on CUDA capable devices of Compute Capability 3.5 and above. This compute

capability was chosen as the capabilities of CUDA Cards differ significantly pre and post Compute Capability 3.

The revision of the initial requirement saved significant time and effort from being wasted in localisation and highly technical activities that would have had very limited benefits for the project.

9.1.1.4 [OA.003] The program shall provide comparative CPU performance metrics.

This requirement was met. The final implementation of the benchmark suite contains both CPU and GPU benchmarks for identical operations. This is then used to provide comparitive performance information to the CAPA reporter in order to make predictions about performance improvements within the codebase.

9.1.1.5 [OA.004] The program shall provide a number of different problem class benchmark algorithms.

This requirement was covered and expanded under sections 9.1.1.1 and 9.1.1.2

9.1.1.6 [OA.005] The program shall provide theoretical performance metrics given a known problem class

This requirement was met. In order to provide the best possible analytics for the serial code analysis, theoretical parallel performance must be understood, so that in situations where a GPU is not present, that relevant calculations may be undertaken to provide an estimate on the anticipated performance. Naturally actual performance and theoretical performance differ significantly for a variety of reasons, however the fundamental considerations involved in algorithm analysis can be known or reasonably estimated from which theoretical performance metrics may be provided. The implementation of this utilised work complexity comparisons between the sequential and parallel versions of operations. Some aspects of the parallel code matched do not provide these theoreticals, as in many cases the final semantics cannot be gathered from the AST representation without runtime information.

9.1.1.7 [OA.006] The program shall include known FPGA performance metrics given a known problem class

This requirement was not met. During the midpoint review of the project it was clear that the original intention of benchmarking CPU, GPU and FPGA implementations of algorithms was an unattainable goal, this was mainly due to the increased focus on the compiler analytics aspect of the projct. Between the original requirements analysis and the midpoint review the emphasis became more focused upon the compiler analytics side, with less emphasis on the relative performances and tradeoffs between the different computing architectures. Due to the size of the project, and the direction it began to proceed, I elected early to not satisfy the requirement, and to remove it from what this project intends to achieve. Removing this requirement provided more time for solving problems which were more relevant to final product.

9.1.1.8 Summary The original design of this project was to primarily provide a benchmarking and comparison suite to assist programmers and engineers in decision making about how to best optimise their software. Very quickly though the project was refocused on analysing source code rather than concepts. This decision increased the challenge of the project, but also provides more utility, as it has the capability to be integrated far deeper in the optimisation decision process. As a result however many aspects of the benchmarking side of the project were compromised. These compromises were recognised before the midpoint of the project and were thoroughly documented in the Design Document.

9.2 Algorithm Analytics Development

This is the crux of my final year project. The core objective was to produce software that analyses a C source file and identify whether the algorithms described may see some benefit from being parallelised. This is extended by the other aspects of this project which in turn provide extra metrics for comparison between CPU and GPU performance. The stretch goal was to provide analysis of an existing C codebase which may contain a variety of potentially parallel algorithms within. As static code analysis is a rather large task to undertake certain decisions have been made to ensure that this project may be completed, and some of these are reflected wtihin the requirements.

9.2.1 Requirements

9.2.1.1 [FR.001] The program shall analyse an algorithm and produce optimisation analysis

This requirement was met. This was the key requirement of the entire project. This requirement could not be compromised on, and as such all other requirements had to relate to ensuring this requirement was met. Analysis of the algorithm was defined as static code analysis, and optimisation analysis was defined as the recognition of potentially parallelizable algorithms within the code. This was achieved through integrating with the Clang tooling, utilising the AST to perform the static analysis. The static analysis itself is primrarly concerned with matching known parallel patterns through the AST Matcher library.

At the midpoint of the project most of this requirement had been met. However I was not pleased with the manner in which the analysis was undertaken. It was quite fragile and difficult to parse. A redevelopment of the entire matching interface was undertaken and as a result I created functional combinator matcher library to facilitate the generation of generic extensible AST Matcher constructs. This allowed for the rapid redevelopment and extension of the original analysers, providing and extensible framework from which further matchers could easily be developed.

The fundamental intention of the requirement was to provide a list of potential improvements from within the code, similar to runtime profiling, by using semantics as described by the designer, rather than performance outcomes determined by a profiler. This was achieved.

9.2.1.2 [FR.002] The program shall produce theoretical performance metrics of developed algorithms

This requirement was met. This relates back to the discussion about theoretical performance metrics in 9.1.1.6. This is the extension of that requirement. Where actual benchmark information is not available the tool provides theoretical performance based on time and work complexities of the relevant algorithms.

9.2.1.3 [FR.004] The source code shall be released under a FOSS license

All source code will be released under a Modified BSD license where permitted.

9.2.1.4 [CA.001] Integrate with the Clang tooling to analyse custom written C code

This requirement has been completely satisfied, a stable build system has been forked from an existing open source project providing the scaffolding around which the entirity of CAPA is built.

9.2.1.5 [CA.002] Identify simple parallel patterns within analysed code

This requirement has been completely satisfied, the three simple parallel patterns for which significant improvements can be found in GPU implementations are:

- Map Operations

- Reductions

- Scans/Prefix Networ/

All three of these simple patterns can be successfully identified within test code.

9.2.1.6 [CA.003] Identify medium complexity parallel patterns within analysed code

This requirement was met. Medium complexity parallel patterns are considered to be patterns within serial code that are clearly parallelizable yet are difficult to identify in a generic sense. This primrarly meant the identification and tagging of 2D Matrix operations. Matrix operations are considered a medium complexity pattern due to the variety of ways in which they may be implemented. As this aspect of the project is primrarly pattern matching and feature detection, identifying the litany of differnet ways a matrix multiplier may be implemented is a significant task. As such writing an accurate, yet generic Matcher and Callback handler for this was a sizeable task. The matrix matcher however was completed, and identifies matrix-matrix multiplication, as well as matrix-vector multiplication.

9.2.1.7 [CA.004] Identify non-trivial parallel patterns within analysed code

Non-trivial parallel patterns mainly defines a broad set of problems that are not clearly parallelizable. The original intention of this requirement was to identify

graph traversals, this however proved to be a particularly difficult task. Whilst graph traversals are not identified by CAPA, we are still able to identify non-trivial potentially parallel sections of code. Primarily CAPA looks for loops with minimal control structures which may be potentially unrolled or vectorised. Additionally CAPA inspects function declarations and types to determine whether they contain types which are to be expected in highly parallel computation, typically pointers, arrays and unsigned integers defining size. This coupled with a low cyclometic complexity within the definition suggests that the included code may be potentially parallelisable, and as such it is often reported by the software.

This requirement has been met, with a slight caveat.

9.2.1.8 [CA.005] Provide Theoretical improvement information

Has been covered extensively here 9.2.1.2 and here 9.1.1.6.

9.2.1.9 [CA.006] Provide more accurate theoretical improvement analysis using additional user specified information

This requirement has been met. Expanding on 9.2.1.2 and 9.1.1.6, if the user decides to provide additional information, then the theoretical performance metrics will take this information into consideration when calculating potential performance improvements. The user is capable of providing information by way of a configuration file which CAPA reads to provide more accurate theoretical calculations.

9.2.1.10 [CA.007] Analyse general C code algorithm descriptions

This requirement has been met. The project is capable of working on any Clang compilable C codebase. As such the analyser is capable of analysing general C code algorithm descriptions from a functional point of view.

9.2.1.11 [CA.008] Analyse existing codebases within tagged regions

This requirement has been conditionally met. Tags are not always visible at every node within the clang AST and as such I was not able to reliably analyse only subsections of a codebase, rather the user has to select which functions within the codebase they do not want to be reported. The program still analyses these regions of the codebase, however upon detecting an ignore clause in the parent function

comment declaration, the analysis for that match terminates and the next matcher begins. This allows the developer to selectively remove false positives or known regions where more computational speed is not necessary. So whilst this condition was originally to only analyse within tagged regions, it has been changed to tag only regions which have not been tagged as: `///CAPA:IGNORE`

9.3   Optimisation Analytics Development

Whilst there are no specific requirements relating to this particular component, this is the unifying feature of the entire project. Combining static code analysis with benchmarks to provide a comprehensive optimisation report, without running a profiler allows for fast identification of potential improvements within a codebase of any size. That is the key intention and aim of this project, and it was achieved.

10 Limitations and Extensions

10.1 Limitations

10.2 Extensions

11 Appendix

All code can be found at:

- `http://github.com/jhana1/CAPA`

- `http://github.com/jhana1/Parallel-AAA`

11.1   Clang AST - Case Study

```
TranslationUnitDecl 0x2efb270 <<invalid sloc>> <invalid sloc>
|-FunctionDecl 0x3055d10 <./CodeUnderTest.cpp:7:1, line:10:1> line:7:6 used random_fill 'void (float *, size_t)'
| |-ParmVarDecl 0x3055b90 <col:18, col:25> col:25 used starting_vec 'float *'
| |-ParmVarDecl 0x3055c00 <col:39, col:46> col:46 used size 'size_t':'unsigned long'
| `-CompoundStmt 0x3056208 <col:51, line:10:1>
|   `-ForStmt 0x30561a0 <line:8:5, line:9:40>
|     |-DeclStmt 0x3055e60 <line:8:10, col:22>
|     | `-VarDecl 0x3055dd0 <col:10, col:21> col:17 used i 'size_t':'unsigned long' cinit
|     |   `-ImplicitCastExpr 0x3055e48 <col:21> 'size_t':'unsigned long' <IntegralCast>
|     |     `-IntegerLiteral 0x3055e28 <col:21> 'int' 0
|     |-<<<NULL>>>
|     |-BinaryOperator 0x3055ef8 <col:24, col:28> '_Bool' '<'
|     | |-ImplicitCastExpr 0x3055ec8 <col:24> 'size_t':'unsigned long' <LValueToRValue>
|     | | `-DeclRefExpr 0x3055e78 <col:24> 'size_t':'unsigned long' lvalue Var 0x3055dd0 'i' 'size_t':'unsigned long'
|     | `-ImplicitCastExpr 0x3055ee0 <col:28> 'size_t':'unsigned long' <LValueToRValue>
|     |   `-DeclRefExpr 0x3055ea0 <col:28> 'size_t':'unsigned long' lvalue ParmVar 0x3055c00 'size' 'size_t':'unsigned long'
|     |-UnaryOperator 0x3055f48 <col:34, col:36> 'size_t':'unsigned long' lvalue prefix '++'
|     | `-DeclRefExpr 0x3055f20 <col:36> 'size_t':'unsigned long' lvalue Var 0x3055dd0 'i' 'size_t':'unsigned long'
|     `-BinaryOperator 0x3056178 <line:9:9, col:40> 'float' lvalue '='
|       |-ArraySubscriptExpr 0x3055fe8 <col:9, col:23> 'float' lvalue
|       | |-ImplicitCastExpr 0x3055fb8 <col:9> 'float *' <LValueToRValue>
|       | | `-DeclRefExpr 0x3055f68 <col:9> 'float *' lvalue ParmVar 0x3055b90 'starting_vec' 'float *'
|       | `-ImplicitCastExpr 0x3055fd0 <col:22> 'size_t':'unsigned long' <LValueToRValue>
|       |   `-DeclRefExpr 0x3055f90 <col:22> 'size_t':'unsigned long' lvalue Var 0x3055dd0 'i' 'size_t':'unsigned long'
|       `-CStyleCastExpr 0x3056150 <col:27, col:40> 'float' <NoOp>
|         `-ImplicitCastExpr 0x3056138 <col:35, col:40> 'float' <IntegralToFloating>
|           `-CallExpr 0x3056100 <col:35, col:40> 'int'
|             `-ImplicitCastExpr 0x30560e8 <col:35> 'int (*)(void) throw()' <FunctionToPointerDecay>
|               `-DeclRefExpr 0x3056068 <col:35> 'int (void) throw()' lvalue Function 0x2fdfcc0 'rand' 'int (void) throw()'
|-FullComment 0x305e640 <line:6:4, col:14>
| `-ParagraphComment 0x305e610 <col:4, col:14>
|   `-TextComment 0x305e5e0 <col:4, col:14> Text="CAPA:IGNORE"
|-FunctionDecl 0x3059950 <line:16:1, line:20:1> line:16:6 used reshape2mat 'void (float *, float **, size_t)'
| |-ParmVarDecl 0x3056240 <col:18, col:25> col:25 used in_vec 'float *'
| |-ParmVarDecl 0x3056350 <col:33, col:48> col:40 used out_vec 'float **':'float **'
| |-ParmVarDecl 0x30563c0 <col:51, col:58> col:58 used dim 'size_t':'unsigned long'
| `-CompoundStmt 0x305a0c8 <col:62, line:20:1>
|   `-ForStmt 0x305a088 <line:17:5, line:19:45>
|     |-DeclStmt 0x3059ab0 <line:17:10, col:22>
|     | `-VarDecl 0x3059a20 <col:10, col:21> col:17 used i 'size_t':'unsigned long' cinit
|     |   `-ImplicitCastExpr 0x3059a98 <col:21> 'size_t':'unsigned long' <IntegralCast>
|     |     `-IntegerLiteral 0x3059a78 <col:21> 'int' 0
|     |-<<<NULL>>>
|     |-BinaryOperator 0x3059b48 <col:24, col:28> '_Bool' '<'
|     | |-ImplicitCastExpr 0x3059b18 <col:24> 'size_t':'unsigned long' <LValueToRValue>
|     | | `-DeclRefExpr 0x3059ac8 <col:24> 'size_t':'unsigned long' lvalue Var 0x3059a20 'i' 'size_t':'unsigned long'
|     | `-ImplicitCastExpr 0x3059b30 <col:28> 'size_t':'unsigned long' <LValueToRValue>
|     |   `-DeclRefExpr 0x3059af0 <col:28> 'size_t':'unsigned long' lvalue ParmVar 0x30563c0 'dim' 'size_t':'unsigned long'
|     |-UnaryOperator 0x3059b98 <col:33, col:35> 'size_t':'unsigned long' lvalue prefix '++'
|     | `-DeclRefExpr 0x3059b70 <col:35> 'size_t':'unsigned long' lvalue Var 0x3059a20 'i' 'size_t':'unsigned long'
|     `-ForStmt 0x305a048 <line:18:9, line:19:45>
|       |-DeclStmt 0x3059c60 <line:18:14, col:26>
|       | `-VarDecl 0x3059bd0 <col:14, col:25> col:21 used j 'size_t':'unsigned long' cinit
|       |   `-ImplicitCastExpr 0x3059c48 <col:25> 'size_t':'unsigned long' <IntegralCast>
|       |     `-IntegerLiteral 0x3059c28 <col:25> 'int' 0
|       |-<<<NULL>>>
|       |-BinaryOperator 0x3059cf8 <col:28, col:32> '_Bool' '<'
|       | |-ImplicitCastExpr 0x3059cc8 <col:28> 'size_t':'unsigned long' <LValueToRValue>
|       | | `-DeclRefExpr 0x3059c78 <col:28> 'size_t':'unsigned long' lvalue Var 0x3059bd0 'j' 'size_t':'unsigned long'
|       | `-ImplicitCastExpr 0x3059ce0 <col:32> 'size_t':'unsigned long' <LValueToRValue>
|       |   `-DeclRefExpr 0x3059ca0 <col:32> 'size_t':'unsigned long' lvalue ParmVar 0x30563c0 'dim' 'size_t':'unsigned long'
|       |-UnaryOperator 0x3059d48 <col:37, col:39> 'size_t':'unsigned long' lvalue prefix '++'
|       | `-DeclRefExpr 0x3059d20 <col:39> 'size_t':'unsigned long' lvalue Var 0x3059bd0 'j' 'size_t':'unsigned long'
|       `-BinaryOperator 0x305a020 <line:19:13, col:45> 'float' lvalue '='
|         |-ArraySubscriptExpr 0x3059e68 <col:13, col:25> 'float' lvalue
|         | |-ImplicitCastExpr 0x3059e38 <col:13, col:22> 'float *' <LValueToRValue>
|         | | `-ArraySubscriptExpr 0x3059de8 <col:13, col:22> 'float *' lvalue
|         | |   |-ImplicitCastExpr 0x3059db8 <col:13> 'float **':'float **' <LValueToRValue>
|         | |   | `-DeclRefExpr 0x3059d68 <col:13> 'float **':'float **' lvalue ParmVar 0x3056350 'out_vec' 'float **':'float **'
|         | |   `-ImplicitCastExpr 0x3059dd0 <col:21> 'size_t':'unsigned long' <LValueToRValue>
|         | |     `-DeclRefExpr 0x3059d90 <col:21> 'size_t':'unsigned long' lvalue Var 0x3059a20 'i' 'size_t':'unsigned long'
|         | `-ImplicitCastExpr 0x3059e50 <col:24> 'size_t':'unsigned long' <LValueToRValue>
|         |   `-DeclRefExpr 0x3059e10 <col:24> 'size_t':'unsigned long' lvalue Var 0x3059bd0 'j' 'size_t':'unsigned long'
|         `-ImplicitCastExpr 0x305a008 <col:29, col:45> 'float' <LValueToRValue>
|           `-ArraySubscriptExpr 0x3059fe0 <col:29, col:45> 'float' lvalue
|             |-ImplicitCastExpr 0x3059fc8 <col:29> 'float *' <LValueToRValue>
|             | `-DeclRefExpr 0x3059e90 <col:29> 'float *' lvalue ParmVar 0x3056240 'in_vec' 'float *'
|             `-BinaryOperator 0x3059fa0 <col:36, col:44> 'unsigned long' '+'
|               |-BinaryOperator 0x3059f38 <col:36, col:38> 'unsigned long' '*'
|               | |-ImplicitCastExpr 0x3059f08 <col:36> 'size_t':'unsigned long' <LValueToRValue>
|               | | `-DeclRefExpr 0x3059eb8 <col:36> 'size_t':'unsigned long' lvalue Var 0x3059a20 'i' 'size_t':'unsigned long'
|               | `-ImplicitCastExpr 0x3059f20 <col:38> 'size_t':'unsigned long' <LValueToRValue>
|               |   `-DeclRefExpr 0x3059ee0 <col:38> 'size_t':'unsigned long' lvalue ParmVar 0x30563c0 'dim' 'size_t':'unsigned long'
|               `-ImplicitCastExpr 0x3059f88 <col:44> 'size_t':'unsigned long' <LValueToRValue>
|                 `-DeclRefExpr 0x3059f60 <col:44> 'size_t':'unsigned long' lvalue Var 0x3059bd0 'j' 'size_t':'unsigned long'
|-FullComment 0x305e760 <line:12:4, line:14:14>
| `-ParagraphComment 0x305e730 <line:12:4, line:14:14>
|   |-TextComment 0x305e6b0 <line:12:4, col:54> Text=" This Function is reponsible for reshaping a vector"
|   |-TextComment 0x305e6d0 <line:13:3, col:17> Text=" into a matrix."
|   `-TextComment 0x305e6f0 <line:14:3, col:14> Text=" CAPA:IGNORE"
|-FunctionDecl 0x305a280 <line:22:1, line:26:1> line:22:6 used reshape2vec 'void (float *, float **, size_t)'
| |-ParmVarDecl 0x305a100 <col:18, col:25> col:25 used out_vec 'float *'
| |-ParmVarDecl 0x305a180 <col:34, col:48> col:41 used in_mat 'float **':'float **'
| |-ParmVarDecl 0x305a1f0 <col:51, col:58> col:58 used dim 'size_t':'unsigned long'
| `-CompoundStmt 0x305aa38 <col:62, line:26:1>
|   `-ForStmt 0x305a9c8 <line:23:5, line:25:45>
|     |-DeclStmt 0x305a3e0 <line:23:10, col:22>
|     | `-VarDecl 0x305a350 <col:10, col:21> col:17 used i 'size_t':'unsigned long' cinit
|     |   `-ImplicitCastExpr 0x305a3c8 <col:21> 'size_t':'unsigned long' <IntegralCast>
|     |     `-IntegerLiteral 0x305a3a8 <col:21> 'int' 0
|     |-<<<NULL>>>
|     |-BinaryOperator 0x305a478 <col:24, col:28> '_Bool' '<'
|     | |-ImplicitCastExpr 0x305a448 <col:24> 'size_t':'unsigned long' <LValueToRValue>
|     | | `-DeclRefExpr 0x305a3f8 <col:24> 'size_t':'unsigned long' lvalue Var 0x305a350 'i' 'size_t':'unsigned long'
|     | `-ImplicitCastExpr 0x305a460 <col:28> 'size_t':'unsigned long' <LValueToRValue>
|     |   `-DeclRefExpr 0x305a420 <col:28> 'size_t':'unsigned long' lvalue ParmVar 0x305a1f0 'dim' 'size_t':'unsigned long'
|     |-UnaryOperator 0x305a4c8 <col:33, col:35> 'size_t':'unsigned long' lvalue prefix '++'
```

```
| |   | `-DeclRefExpr 0x305a4a0 <col:35> 'size_t':'unsigned long' lvalue Var 0x305a350 'i' 'size_t':'unsigned long'
| `-ForStmt 0x305a988 <line:24:9, line:25:45>
|   |-DeclStmt 0x305a590 <line:24:14, col:26>
|   | `-VarDecl 0x305a500 <col:14, col:25> col:21 used j 'size_t':'unsigned long' cinit
|   |   `-ImplicitCastExpr 0x305a578 <col:25> 'size_t':'unsigned long' <IntegralCast>
|   |     `-IntegerLiteral 0x305a558 <col:25> 'int' 0
|   |-<<<NULL>>>
|   |-BinaryOperator 0x305a628 <col:28, col:32> '_Bool' '<'
|   | |-ImplicitCastExpr 0x305a5f8 <col:28> 'size_t':'unsigned long' <LValueToRValue>
|   | | `-DeclRefExpr 0x305a5a8 <col:28> 'size_t':'unsigned long' lvalue Var 0x305a500 'j' 'size_t':'unsigned long'
|   | `-ImplicitCastExpr 0x305a610 <col:32> 'size_t':'unsigned long' <LValueToRValue>
|   |   `-DeclRefExpr 0x305a5d0 <col:32> 'size_t':'unsigned long' lvalue ParmVar 0x305a1f0 'dim' 'size_t':'unsigned long'
|   |-UnaryOperator 0x305a678 <col:37, col:39> 'size_t':'unsigned long' lvalue prefix '++'
|   | `-DeclRefExpr 0x305a650 <col:39> 'size_t':'unsigned long' lvalue Var 0x305a500 'j' 'size_t':'unsigned long'
|   `-BinaryOperator 0x305a960 <line:25:13, col:45> 'float' lvalue '='
|     |-ArraySubscriptExpr 0x305a7e8 <col:13, col:30> 'float' lvalue
|     | |-ImplicitCastExpr 0x305a7d0 <col:13> 'float *' <LValueToRValue>
|     | | `-DeclRefExpr 0x305a698 <col:13> 'float *' lvalue ParmVar 0x305a100 'out_vec' 'float *'
|     | `-BinaryOperator 0x305a7a8 <col:21, col:29> 'unsigned long' '+'
|     |   |-BinaryOperator 0x305a740 <col:21, col:23> 'unsigned long' '*'
|     |   | |-ImplicitCastExpr 0x305a710 <col:21> 'size_t':'unsigned long' <LValueToRValue>
|     |   | | `-DeclRefExpr 0x305a6c0 <col:21> 'size_t':'unsigned long' lvalue Var 0x305a350 'i' 'size_t':'unsigned long'
|     |   | `-ImplicitCastExpr 0x305a728 <col:23> 'size_t':'unsigned long' <LValueToRValue>
|     |   |   `-DeclRefExpr 0x305a6e8 <col:23> 'size_t':'unsigned long' lvalue ParmVar 0x305a1f0 'dim' 'size_t':'unsigned long'
|     |   `-ImplicitCastExpr 0x305a790 <col:29> 'size_t':'unsigned long' <LValueToRValue>
|     |     `-DeclRefExpr 0x305a768 <col:29> 'size_t':'unsigned long' lvalue Var 0x305a500 'j' 'size_t':'unsigned long'
|     `-ImplicitCastExpr 0x305a948 <col:34, col:45> 'float' <LValueToRValue>
|       `-ArraySubscriptExpr 0x305a920 <col:34, col:45> 'float' lvalue
|         |-ImplicitCastExpr 0x305a8e0 <col:34, col:42> 'float *' <LValueToRValue>
|         | `-ArraySubscriptExpr 0x305a890 <col:34, col:42> 'float *' lvalue
|         |   |-ImplicitCastExpr 0x305a860 <col:34> 'float **':'float **' <LValueToRValue>
|         |   | `-DeclRefExpr 0x305a810 <col:34> 'float **':'float **' lvalue ParmVar 0x305a180 'in_mat' 'float **':'float **'
|         |   `-ImplicitCastExpr 0x305a878 <col:41> 'size_t':'unsigned long' <LValueToRValue>
|         |     `-DeclRefExpr 0x305a838 <col:41> 'size_t':'unsigned long' lvalue Var 0x305a350 'i' 'size_t':'unsigned long'
|         `-ImplicitCastExpr 0x305a8f8 <col:44> 'size_t':'unsigned long' <LValueToRValue>
|           `-DeclRefExpr 0x305a8b8 <col:44> 'size_t':'unsigned long' lvalue Var 0x305a500 'j' 'size_t':'unsigned long'
|-FunctionDecl 0x305ad00 <line:29:1, col:55> col:6 used mmult 'void (float **, float **, float **, size_t)'
| |-ParmVarDecl 0x305aa70 <col:12, col:20> col:20 A 'float **'
| |-ParmVarDecl 0x305aae0 <col:23, col:31> col:31 B 'float **'
| |-ParmVarDecl 0x305ab50 <col:34, col:42> col:42 C 'float **'
| |-ParmVarDecl 0x305abc0 <col:45, col:52> col:52 dim 'size_t':'unsigned long'
| `-FullComment 0x305e830 <line:28:4, col:15>
|   `-ParagraphComment 0x305e800 <col:4, col:15>
|     `-TextComment 0x305e7d0 <col:4, col:15> Text=" CAPA:IGNORE"
`-FunctionDecl 0x305ade0 <line:31:1, line:74:1> line:31:5 main 'int (void)'
  `-CompoundStmt 0x305d810 <line:32:1, line:74:1>
    |-DeclStmt 0x305afb8 <line:34:5, col:35>
    | `-VarDecl 0x305ae90 <col:5, col:31> col:18 referenced ELEMS 'const size_t':'const unsigned long' cinit
    |   `-ImplicitCastExpr 0x305af50 <col:26, col:31> 'const size_t':'const unsigned long' <IntegralCast>
    |     `-BinaryOperator 0x305af28 <col:26, col:31> 'int' '*'
    |       |-IntegerLiteral 0x305aee8 <col:26> 'int' 1000
    |       `-IntegerLiteral 0x305af08 <col:31> 'int' 1000
    |-DeclStmt 0x305b0c8 <line:35:5, col:30>
    | `-VarDecl 0x305b070 <col:5, col:29> col:11 used starting_vec 'float [1000000]'
    |-DeclStmt 0x305b148 <line:36:5, col:13>
    | `-VarDecl 0x305b0f0 <col:5, col:12> col:12 used t 'time_t':'long'
    |-CallExpr 0x305b410 <line:38:5, col:30> 'void'
    | |-ImplicitCastExpr 0x305b3f8 <col:5> 'void (*)(unsigned int) throw()' <FunctionToPointerDecay>
    | | `-DeclRefExpr 0x305b370 <col:5> 'void (unsigned int) throw()' lvalue Function 0x2fdfdf0 'srand' 'void (unsigned int) throw()'
    | `-CStyleCastExpr 0x305b348 <col:11, col:29> 'unsigned int' <NoOp>
    |   `-ImplicitCastExpr 0x305b330 <col:22, col:29> 'unsigned int' <IntegralCast>
    |     `-CallExpr 0x305b2f0 <col:22, col:29> 'time_t':'long'
    |       |-ImplicitCastExpr 0x305b2d8 <col:22> 'time_t (*)(time_t *) throw()' <FunctionToPointerDecay>
    |       | `-DeclRefExpr 0x305b258 <col:22> 'time_t (time_t *) throw()' lvalue Function 0x3001960 'time' 'time_t (time_t *) throw()'
    |       `-UnaryOperator 0x305b238 <col:27, col:28> 'time_t *' prefix '&'
    |         `-DeclRefExpr 0x305b210 <col:28> 'time_t':'long' lvalue Var 0x305b0f0 't' 'time_t':'long'
    |-CallExpr 0x305b580 <line:39:5, col:36> 'void'
    | |-ImplicitCastExpr 0x305b568 <col:5> 'void (*)(float *, size_t)' <FunctionToPointerDecay>
    | | `-DeclRefExpr 0x305b4e8 <col:5> 'void (float *, size_t)' lvalue Function 0x3055d10 'random_fill' 'void (float *, size_t)'
    | |-ImplicitCastExpr 0x305b5b8 <col:17> 'float *' <ArrayToPointerDecay>
    | | `-DeclRefExpr 0x305b498 <col:17> 'float [1000000]' lvalue Var 0x305b070 'starting_vec' 'float [1000000]'
    | `-ImplicitCastExpr 0x305b5d0 <col:31> 'size_t':'unsigned long' <LValueToRValue>
    |   `-DeclRefExpr 0x305b4c0 <col:31> 'const size_t':'const unsigned long' lvalue Var 0x305ae90 'ELEMS' 'const size_t':'const unsigned long'
    |-ForStmt 0x305ba08 <line:42:5, line:45:5>
    | |-DeclStmt 0x305b690 <line:42:10, col:22>
    | | `-VarDecl 0x305b600 <col:10, col:21> col:17 used i 'size_t':'unsigned long' cinit
    | |   `-ImplicitCastExpr 0x305b678 <col:21> 'size_t':'unsigned long' <IntegralCast>
    | |     `-IntegerLiteral 0x305b658 <col:21> 'int' 0
    | |-<<<NULL>>>
    | |-BinaryOperator 0x305b728 <col:24, col:28> '_Bool' '<'
    | | |-ImplicitCastExpr 0x305b6f8 <col:24> 'size_t':'unsigned long' <LValueToRValue>
    | | | `-DeclRefExpr 0x305b6a8 <col:24> 'size_t':'unsigned long' lvalue Var 0x305b600 'i' 'size_t':'unsigned long'
    | | `-ImplicitCastExpr 0x305b710 <col:28> 'size_t':'unsigned long' <LValueToRValue>
    | |   `-DeclRefExpr 0x305b6d0 <col:28> 'const size_t':'const unsigned long' lvalue Var 0x305ae90 'ELEMS' 'const size_t':'const unsigned long'
    | |-UnaryOperator 0x305b778 <col:35, col:37> 'size_t':'unsigned long' lvalue prefix '++'
    | | `-DeclRefExpr 0x305b750 <col:37> 'size_t':'unsigned long' lvalue Var 0x305b600 'i' 'size_t':'unsigned long'
    | `-CompoundStmt 0x305b9e0 <col:39, line:45:5>
    |   |-CompoundAssignOperator 0x305b878 <line:43:9, col:28> 'float' lvalue '/=' ComputeLHSTy='float' ComputeResultTy='float'
    |   | |-ArraySubscriptExpr 0x305b818 <col:9, col:23> 'float' lvalue
    |   | | |-ImplicitCastExpr 0x305b7e8 <col:9> 'float *' <ArrayToPointerDecay>
    |   | | | `-DeclRefExpr 0x305b798 <col:9> 'float [1000000]' lvalue Var 0x305b070 'starting_vec' 'float [1000000]'
    |   | | `-ImplicitCastExpr 0x305b800 <col:22> 'size_t':'unsigned long' <LValueToRValue>
    |   | |   `-DeclRefExpr 0x305b7c0 <col:22> 'size_t':'unsigned long' lvalue Var 0x305b600 'i' 'size_t':'unsigned long'
    |   | `-ImplicitCastExpr 0x305b860 <col:28> 'float' <IntegralToFloating>
    |   |   `-IntegerLiteral 0x305b840 <col:28> 'int' 2
    |   `-CompoundAssignOperator 0x305b9a8 <line:44:9, col:28> 'float' lvalue '+=' ComputeLHSTy='float' ComputeResultTy='float'
    |     |-ArraySubscriptExpr 0x305b948 <col:9, col:23> 'float' lvalue
    |     | |-ImplicitCastExpr 0x305b900 <col:9> 'float *' <ArrayToPointerDecay>
    |     | | `-DeclRefExpr 0x305b8b0 <col:9> 'float [1000000]' lvalue Var 0x305b070 'starting_vec' 'float [1000000]'
    |     | `-ImplicitCastExpr 0x305b930 <col:22> 'size_t':'unsigned long' <LValueToRValue>
    |     |   `-DeclRefExpr 0x305b8d8 <col:22> 'size_t':'unsigned long' lvalue Var 0x305b600 'i' 'size_t':'unsigned long'
    |     `-ImplicitCastExpr 0x305b990 <col:28> 'float' <IntegralToFloating>
    |       `-IntegerLiteral 0x305b970 <col:28> 'int' 4
    |-DeclStmt 0x305baf0 <line:48:5, col:16>
    | `-VarDecl 0x305ba60 <col:5, col:15> col:11 used k 'float' cinit
    |   `-ImplicitCastExpr 0x305bad8 <col:15> 'float' <IntegralToFloating>
    |     `-IntegerLiteral 0x305bab8 <col:15> 'int' 0
```

```
ForStmt 0x305be58 <line:49:5, line:50:30>
| DeclStmt 0x305bbb0 <line:49:10, col:22>
| `-VarDecl 0x305bb20 <col:10, col:21> col:17 used i 'size_t':'unsigned long' cinit
|   `-ImplicitCastExpr 0x305bb98 <col:21> 'size_t':'unsigned long' <IntegralCast>
|     `-IntegerLiteral 0x305bb78 <col:21> 'int' 0
| <<<NULL>>>
| BinaryOperator 0x305bc48 <col:24, col:28> '_Bool' '<'
| |-ImplicitCastExpr 0x305bc18 <col:24> 'size_t':'unsigned long' <LValueToRValue>
| | `-DeclRefExpr 0x305bbc8 <col:24> 'size_t':'unsigned long' lvalue Var 0x305bb20 'i' 'size_t':'unsigned long'
| `-ImplicitCastExpr 0x305bc30 <col:28> 'size_t':'unsigned long' <LValueToRValue>
|   `-DeclRefExpr 0x305bbf0 <col:28> 'const size_t':'const unsigned long' lvalue Var 0x305ae90 'ELEMS' 'const size_t':'const unsigned long'
| UnaryOperator 0x305bc98 <col:35, col:37> 'size_t':'unsigned long' lvalue prefix '++'
| `-DeclRefExpr 0x305bc70 <col:37> 'size_t':'unsigned long' lvalue Var 0x305bb20 'i' 'size_t':'unsigned long'
`-CompoundAssignOperator 0x305be20 <line:50:9, col:30> 'float' lvalue '+=' ComputeLHSTy='float' ComputeResultTy='float'
  |-DeclRefExpr 0x305bcb8 <col:9> 'float' lvalue Var 0x305ba60 'k' 'float'
  `-BinaryOperator 0x305bdf8 <col:14, col:30> 'float' '/'
    |-ImplicitCastExpr 0x305bdb0 <col:14, col:28> 'float' <LValueToRValue>
    | `-ArraySubscriptExpr 0x305bd60 <col:14, col:28> 'float' lvalue
    |   |-ImplicitCastExpr 0x305bd30 <col:14> 'float *' <ArrayToPointerDecay>
    |   | `-DeclRefExpr 0x305bce0 <col:14> 'float [1000000]' lvalue Var 0x305b070 'starting_vec' 'float [1000000]'
    |   `-ImplicitCastExpr 0x305bd48 <col:27> 'size_t':'unsigned long' <LValueToRValue>
    |     `-DeclRefExpr 0x305bd08 <col:27> 'size_t':'unsigned long' lvalue Var 0x305bb20 'i' 'size_t':'unsigned long'
    `-ImplicitCastExpr 0x305bde0 <col:30> 'float' <IntegralToFloating>
      `-ImplicitCastExpr 0x305bdc8 <col:30> 'size_t':'unsigned long' <LValueToRValue>
        `-DeclRefExpr 0x305bd88 <col:30> 'const size_t':'const unsigned long' lvalue Var 0x305ae90 'ELEMS' 'const size_t':'const unsigned long'
DeclStmt 0x305bf58 <line:53:5, col:25>
`-VarDecl 0x305bf00 <col:5, col:24> col:11 used cum_sum 'float [1000000]'
BinaryOperator 0x305c100 <line:54:5, col:34> 'float' lvalue '='
|-ArraySubscriptExpr 0x305bfd0 <col:5, col:14> 'float' lvalue
| |-ImplicitCastExpr 0x305bfb8 <col:5> 'float *' <ArrayToPointerDecay>
| | `-DeclRefExpr 0x305bf70 <col:5> 'float [1000000]' lvalue Var 0x305bf00 'cum_sum' 'float [1000000]'
| `-IntegerLiteral 0x305bf98 <col:13> 'int' 0
`-BinaryOperator 0x305c0d8 <col:18, col:34> 'float' '/'
  |-ImplicitCastExpr 0x305c0a8 <col:18, col:32> 'float' <LValueToRValue>
  | `-ArraySubscriptExpr 0x305c058 <col:18, col:32> 'float' lvalue
  |   |-ImplicitCastExpr 0x305c040 <col:18> 'float *' <ArrayToPointerDecay>
  |   | `-DeclRefExpr 0x305bff8 <col:18> 'float [1000000]' lvalue Var 0x305b070 'starting_vec' 'float [1000000]'
  |   `-IntegerLiteral 0x305c020 <col:31> 'int' 0
  `-ImplicitCastExpr 0x305c0c0 <col:34> 'float' <LValueToRValue>
    `-DeclRefExpr 0x305c080 <col:34> 'float' lvalue Var 0x305ba60 'k' 'float'
ForStmt 0x305c630 <line:55:5, line:56:57>
| DeclStmt 0x305c1d0 <line:55:10, col:22>
| `-VarDecl 0x305c140 <col:10, col:21> col:17 used i 'size_t':'unsigned long' cinit
|   `-ImplicitCastExpr 0x305c1b8 <col:21> 'size_t':'unsigned long' <IntegralCast>
|     `-IntegerLiteral 0x305c198 <col:21> 'int' 1
| <<<NULL>>>
| BinaryOperator 0x305c268 <col:24, col:28> '_Bool' '<'
| |-ImplicitCastExpr 0x305c238 <col:24> 'size_t':'unsigned long' <LValueToRValue>
| | `-DeclRefExpr 0x305c1e8 <col:24> 'size_t':'unsigned long' lvalue Var 0x305c140 'i' 'size_t':'unsigned long'
| `-ImplicitCastExpr 0x305c250 <col:28> 'size_t':'unsigned long' <LValueToRValue>
|   `-DeclRefExpr 0x305c210 <col:28> 'const size_t':'const unsigned long' lvalue Var 0x305ae90 'ELEMS' 'const size_t':'const unsigned long'
| UnaryOperator 0x305c2b8 <col:35, col:36> 'size_t':'unsigned long' postfix '++'
| `-DeclRefExpr 0x305c290 <col:35> 'size_t':'unsigned long' lvalue Var 0x305c140 'i' 'size_t':'unsigned long'
`-BinaryOperator 0x305c608 <line:56:9, col:57> 'float' lvalue '='
  |-ArraySubscriptExpr 0x305c358 <col:9, col:18> 'float' lvalue
  | |-ImplicitCastExpr 0x305c328 <col:9> 'float *' <ArrayToPointerDecay>
  | | `-DeclRefExpr 0x305c2d8 <col:9> 'float [1000000]' lvalue Var 0x305bf00 'cum_sum' 'float [1000000]'
  | `-ImplicitCastExpr 0x305c340 <col:17> 'size_t':'unsigned long' <LValueToRValue>
  |   `-DeclRefExpr 0x305c300 <col:17> 'size_t':'unsigned long' lvalue Var 0x305c140 'i' 'size_t':'unsigned long'
  `-BinaryOperator 0x305c5e0 <col:22, col:57> 'float' '+'
    |-BinaryOperator 0x305c498 <col:22, col:38> 'float' '/'
    | |-ImplicitCastExpr 0x305c450 <col:22, col:36> 'float' <LValueToRValue>
    | | `-ArraySubscriptExpr 0x305c400 <col:22, col:36> 'float' lvalue
    | |   |-ImplicitCastExpr 0x305c3d0 <col:22> 'float *' <ArrayToPointerDecay>
    | |   | `-DeclRefExpr 0x305c380 <col:22> 'float [1000000]' lvalue Var 0x305b070 'starting_vec' 'float [1000000]'
    | |   `-ImplicitCastExpr 0x305c3e8 <col:35> 'size_t':'unsigned long' <LValueToRValue>
    | |     `-DeclRefExpr 0x305c3a8 <col:35> 'size_t':'unsigned long' lvalue Var 0x305c140 'i' 'size_t':'unsigned long'
    | `-ImplicitCastExpr 0x305c480 <col:38> 'float' <IntegralToFloating>
    |   `-ImplicitCastExpr 0x305c468 <col:38> 'size_t':'unsigned long' <LValueToRValue>
    |     `-DeclRefExpr 0x305c428 <col:38> 'const size_t':'const unsigned long' lvalue Var 0x305ae90 'ELEMS' 'const size_t':'const unsigned long'
    `-ImplicitCastExpr 0x305c5c8 <col:46, col:57> 'float' <LValueToRValue>
      `-ArraySubscriptExpr 0x305c5a0 <col:46, col:57> 'float' lvalue
        |-ImplicitCastExpr 0x305c588 <col:46> 'float *' <ArrayToPointerDecay>
        | `-DeclRefExpr 0x305c4c0 <col:46> 'float [1000000]' lvalue Var 0x305bf00 'cum_sum' 'float [1000000]'
        `-BinaryOperator 0x305c560 <col:54, col:56> 'unsigned long' '-'
          |-ImplicitCastExpr 0x305c530 <col:54> 'size_t':'unsigned long' <LValueToRValue>
          | `-DeclRefExpr 0x305c4e8 <col:54> 'size_t':'unsigned long' lvalue Var 0x305c140 'i' 'size_t':'unsigned long'
          `-ImplicitCastExpr 0x305c548 <col:56> 'unsigned long' <IntegralCast>
            `-IntegerLiteral 0x305c510 <col:56> 'int' 1
DeclStmt 0x305c8b8 <line:59:5, col:35>
`-VarDecl 0x305c680 <col:5, col:34> col:18 used dim 'const size_t':'const unsigned long' cinit
  `-ImplicitCastExpr 0x305c850 <col:24, col:34> 'const size_t':'const unsigned long' <FloatingToIntegral>
    `-CallExpr 0x305c7f0 <col:24, col:34> 'double'
      |-ImplicitCastExpr 0x305c7d8 <col:24> 'double (*)(double) throw()' <FunctionToPointerDecay>
      | `-DeclRefExpr 0x305c758 <col:24> 'double (double) throw()' lvalue Function 0x2fc3370 'sqrt' 'double (double) throw()'
      `-ImplicitCastExpr 0x305c838 <col:29> 'double' <IntegralToFloating>
        `-ImplicitCastExpr 0x305c820 <col:29> 'size_t':'unsigned long' <LValueToRValue>
          `-DeclRefExpr 0x305c730 <col:29> 'const size_t':'const unsigned long' lvalue Var 0x305ae90 'ELEMS' 'const size_t':'const unsigned long'
DeclStmt 0x305ca78 <line:60:5, col:32>
`-VarDecl 0x305ca20 <col:5, col:31> col:11 used cum_sum_mat 'float [dim][dim]'
CallExpr 0x305cc80 <line:62:5, col:53> 'void'
|-ImplicitCastExpr 0x305cc68 <col:5> 'void (*)(float *, float **, size_t)' <FunctionToPointerDecay>
| `-DeclRefExpr 0x305cbe0 <col:5> 'void (float *, float **, size_t)' lvalue Function 0x3059950 'reshape2mat' 'void (float *, float **, size_t)'
|-ImplicitCastExpr 0x305ccc0 <col:17> 'float *' <ArrayToPointerDecay>
| `-DeclRefExpr 0x305cae8 <col:17> 'float [1000000]' lvalue Var 0x305bf00 'cum_sum' 'float [1000000]'
|-CStyleCastExpr 0x305cb90 <col:26, col:37> 'float **' <BitCast>
| `-ImplicitCastExpr 0x305cb78 <col:37> 'float (*)[dim]' <ArrayToPointerDecay>
|   `-DeclRefExpr 0x305cb10 <col:37> 'float [dim][dim]' lvalue Var 0x305ca20 'cum_sum_mat' 'float [dim][dim]'
`-ImplicitCastExpr 0x305ccd8 <col:50> 'size_t':'unsigned long' <LValueToRValue>
  `-DeclRefExpr 0x305cbb8 <col:50> 'const size_t':'const unsigned long' lvalue Var 0x305c680 'dim' 'const size_t':'const unsigned long'
CallExpr 0x305cf90 <line:63:5, col:86> 'void'
|-ImplicitCastExpr 0x305cf78 <col:5> 'void (*)(float **, float **, float **, size_t)' <FunctionToPointerDecay>
| `-DeclRefExpr 0x305cef0 <col:5> 'void (float **, float **, float **, size_t)' lvalue Function 0x305ad00 'mmult' 'void (float **, float **, float **, size_t)'
|-CStyleCastExpr 0x305cda0 <col:11, col:22> 'float **' <BitCast>
| `-ImplicitCastExpr 0x305cd88 <col:22> 'float (*)[dim]' <ArrayToPointerDecay>
|   `-DeclRefExpr 0x305cd48 <col:22> 'float [dim][dim]' lvalue Var 0x305ca20 'cum_sum_mat' 'float [dim][dim]'
|-CStyleCastExpr 0x305ce20 <col:35, col:46> 'float **' <BitCast>
```

```
|   |   |   ImplicitCastExpr 0x305ce08 <col:46> 'float (*)[dim]' <ArrayToPointerDecay>
|   |   |     DeclRefExpr 0x305cdc8 <col:46> 'float [dim][dim]' lvalue Var 0x305ca20 'cum_sum_mat' 'float [dim][dim]'
|   |   | CStyleCastExpr 0x305cea0 <col:59, col:70> 'float **' <BitCast>
|   |   |   ImplicitCastExpr 0x305ce88 <col:70> 'float (*)[dim]' <ArrayToPointerDecay>
|   |   |     DeclRefExpr 0x305ce48 <col:70> 'float [dim][dim]' lvalue Var 0x305ca20 'cum_sum_mat' 'float [dim][dim]'
|   |   | ImplicitCastExpr 0x305cfd8 <col:83> 'size_t':'unsigned long' <LValueToRValue>
|   |   |   DeclRefExpr 0x305cec8 <col:83> 'const size_t':'const unsigned long' lvalue Var 0x305c680 'dim' 'const size_t':'const unsigned long'
|   | CallExpr 0x305d158 <line:64:5, col:53> 'void'
|   | | ImplicitCastExpr 0x305d140 <col:5> 'void (*)(float *, float **, size_t)' <FunctionToPointerDecay>
|   | |   DeclRefExpr 0x305d118 <col:5> 'void (float *, float **, size_t)' lvalue Function 0x305a280 'reshape2vec' 'void (float *, float **, size_t)'
|   | | ImplicitCastExpr 0x305d198 <col:17> 'float *' <ArrayToPointerDecay>
|   | |   DeclRefExpr 0x305d048 <col:17> 'float [1000000]' lvalue Var 0x305bf00 'cum_sum' 'float [1000000]'
|   | | CStyleCastExpr 0x305d0c8 <col:26, col:37> 'float **' <BitCast>
|   | |   ImplicitCastExpr 0x305d0b0 <col:37> 'float (*)[dim]' <ArrayToPointerDecay>
|   | |     DeclRefExpr 0x305d070 <col:37> 'float [dim][dim]' lvalue Var 0x305ca20 'cum_sum_mat' 'float [dim][dim]'
|   | | ImplicitCastExpr 0x305d1b0 <col:50> 'size_t':'unsigned long' <LValueToRValue>
|   | |   DeclRefExpr 0x305d0f0 <col:50> 'const size_t':'const unsigned long' lvalue Var 0x305c680 'dim' 'const size_t':'const unsigned long'
|   | BinaryOperator 0x305d228 <line:68:5, col:9> 'float' lvalue '='
|   | | DeclRefExpr 0x305d1c8 <col:5> 'float' lvalue Var 0x305ba60 'k' 'float'
|   | | ImplicitCastExpr 0x305d210 <col:9> 'float' <IntegralToFloating>
|   | |   IntegerLiteral 0x305d1f0 <col:9> 'int' 0
|   | ForStmt 0x305d758 <line:69:5, line:71:38>
|   | | DeclStmt 0x305d2f0 <line:69:10, col:22>
|   | | | VarDecl 0x305d260 <col:10, col:21> col:17 used i 'size_t':'unsigned long' cinit
|   | | |   ImplicitCastExpr 0x305d2d8 <col:21> 'size_t':'unsigned long' <IntegralCast>
|   | | |     IntegerLiteral 0x305d2b8 <col:21> 'int' 0
|   | | |-<<<NULL>>>
|   | | BinaryOperator 0x305d388 <col:24, col:28> '_Bool' '<'
|   | | | ImplicitCastExpr 0x305d358 <col:24> 'size_t':'unsigned long' <LValueToRValue>
|   | | |   DeclRefExpr 0x305d308 <col:24> 'size_t':'unsigned long' lvalue Var 0x305d260 'i' 'size_t':'unsigned long'
|   | | | ImplicitCastExpr 0x305d370 <col:28> 'size_t':'unsigned long' <LValueToRValue>
|   | | |   DeclRefExpr 0x305d330 <col:28> 'const size_t':'const unsigned long' lvalue Var 0x305ae90 'ELEMS' 'const size_t':'const unsigned long'
|   | | UnaryOperator 0x305d3d8 <col:35, col:37> 'size_t':'unsigned long' lvalue prefix '++'
|   | | | DeclRefExpr 0x305d3b0 <col:37> 'size_t':'unsigned long' lvalue Var 0x305d260 'i' 'size_t':'unsigned long'
|   | | IfStmt 0x305d728 <line:70:9, line:71:38>
|   | | |-<<<NULL>>>
|   | | | UnaryOperator 0x305d4d0 <line:70:13, col:20> '_Bool' prefix '!'
|   | | | | ImplicitCastExpr 0x305d4b8 <col:14, col:20> '_Bool' <IntegralToBoolean>
|   | | | |   ParenExpr 0x305d498 <col:14, col:20> 'unsigned long'
|   | | | |     BinaryOperator 0x305d470 <col:15, col:19> 'unsigned long' '%'
|   | | | |     | ImplicitCastExpr 0x305d440 <col:15> 'size_t':'unsigned long' <LValueToRValue>
|   | | | |     |   DeclRefExpr 0x305d3f8 <col:15> 'size_t':'unsigned long' lvalue Var 0x305d260 'i' 'size_t':'unsigned long'
|   | | | |     | ImplicitCastExpr 0x305d458 <col:19> 'unsigned long' <IntegralCast>
|   | | | |     |   IntegerLiteral 0x305d420 <col:19> 'int' 2
|   | | | BinaryOperator 0x305d700 <line:71:13, col:38> 'float' lvalue '='
|   | | | | DeclRefExpr 0x305d4f0 <col:13> 'float' lvalue Var 0x305ba60 'k' 'float'
|   | | | | BinaryOperator 0x305d6d8 <col:17, col:38> 'float' '+'
|   | | | | | ImplicitCastExpr 0x305d6a8 <col:17> 'float' <LValueToRValue>
|   | | | | |   DeclRefExpr 0x305d518 <col:17> 'float' lvalue Var 0x305ba60 'k' 'float'
|   | | | | | ImplicitCastExpr 0x305d6c0 <col:21, col:38> 'float' <LValueToRValue>
|   | | | | |   ArraySubscriptExpr 0x305d680 <col:21, col:38> 'float' lvalue
|   | | | | |     ImplicitCastExpr 0x305d668 <col:21> 'float *' <ArrayToPointerDecay>
|   | | | | |     | DeclRefExpr 0x305d540 <col:21> 'float [1000000]' lvalue Var 0x305bf00 'cum_sum' 'float [1000000]'
|   | | | | |     BinaryOperator 0x305d640 <col:29, col:37> 'unsigned long' '-'
|   | | | | |     | BinaryOperator 0x305d5e0 <col:29, col:33> 'unsigned long' '+'
|   | | | | |     | | ImplicitCastExpr 0x305d5b0 <col:29> 'size_t':'unsigned long' <LValueToRValue>
|   | | | | |     | |   DeclRefExpr 0x305d568 <col:29> 'size_t':'unsigned long' lvalue Var 0x305d260 'i' 'size_t':'unsigned long'
|   | | | | |     | | ImplicitCastExpr 0x305d5c8 <col:33> 'unsigned long' <IntegralCast>
|   | | | | |     | |   IntegerLiteral 0x305d590 <col:33> 'int' 1
|   | | | | |     | ImplicitCastExpr 0x305d628 <col:37> 'unsigned long' <IntegralCast>
|   | | | | |     |   IntegerLiteral 0x305d608 <col:37> 'int' 1
|   | | | |-<<<NULL>>>
|   | ReturnStmt 0x305d7f0 <line:73:5, col:12>
|   |   ImplicitCastExpr 0x305d7d8 <col:12> 'int' <FloatingToIntegral>
|   |     ImplicitCastExpr 0x305d7c0 <col:12> 'float' <LValueToRValue>
|   |       DeclRefExpr 0x305d798 <col:12> 'float' lvalue Var 0x305ba60 'k' 'float'
|-FunctionDecl 0x305daf0 prev 0x305ad00 <line:76:1, line:81:1> line:76:6 used mmult 'void (float **, float **, float **, size_t)'
| |-ParmVarDecl 0x305d8e0 <col:12, col:20> col:20 used A 'float **'
| |-ParmVarDecl 0x305d970 <col:23, col:31> col:31 used B 'float **'
| |-ParmVarDecl 0x305d9e0 <col:34, col:42> col:42 used C 'float **'
| |-ParmVarDecl 0x305da50 <col:45, col:52> col:52 used dim 'size_t':'unsigned long'
| `-CompoundStmt 0x305e580 <col:56, line:81:1>
|   `-ForStmt 0x305e540 <line:77:5, line:80:44>
|     |-DeclStmt 0x305dc50 <line:77:10, col:22>
|     | `-VarDecl 0x305dbc0 <col:10, col:21> col:17 used i 'size_t':'unsigned long' cinit
|     |   `-ImplicitCastExpr 0x305dc38 <col:21> 'size_t':'unsigned long' <IntegralCast>
|     |     `-IntegerLiteral 0x305dc18 <col:21> 'int' 0
|     |-<<<NULL>>>
|     |-BinaryOperator 0x305dce8 <col:24, col:28> '_Bool' '<'
|     | |-ImplicitCastExpr 0x305dcb8 <col:24> 'size_t':'unsigned long' <LValueToRValue>
|     | | `-DeclRefExpr 0x305dc68 <col:24> 'size_t':'unsigned long' lvalue Var 0x305dbc0 'i' 'size_t':'unsigned long'
|     | `-ImplicitCastExpr 0x305dcd0 <col:28> 'size_t':'unsigned long' <LValueToRValue>
|     |   `-DeclRefExpr 0x305dc90 <col:28> 'size_t':'unsigned long' lvalue ParmVar 0x305da50 'dim' 'size_t':'unsigned long'
|     |-UnaryOperator 0x305dd38 <col:33, col:35> 'size_t':'unsigned long' lvalue prefix '++'
|     | `-DeclRefExpr 0x305dd10 <col:35> 'size_t':'unsigned long' lvalue Var 0x305dbc0 'i' 'size_t':'unsigned long'
|     `-ForStmt 0x305e500 <line:78:9, line:80:44>
|       |-DeclStmt 0x305de00 <line:78:14, col:26>
|       | `-VarDecl 0x305dd70 <col:14, col:25> col:21 used j 'size_t':'unsigned long' cinit
|       |   `-ImplicitCastExpr 0x305dde8 <col:25> 'size_t':'unsigned long' <IntegralCast>
|       |     `-IntegerLiteral 0x305ddc8 <col:25> 'int' 0
|       |-<<<NULL>>>
|       |-BinaryOperator 0x305de98 <col:28, col:32> '_Bool' '<'
|       | |-ImplicitCastExpr 0x305de68 <col:28> 'size_t':'unsigned long' <LValueToRValue>
|       | | `-DeclRefExpr 0x305de18 <col:28> 'size_t':'unsigned long' lvalue Var 0x305dd70 'j' 'size_t':'unsigned long'
|       | `-ImplicitCastExpr 0x305de80 <col:32> 'size_t':'unsigned long' <LValueToRValue>
|       |   `-DeclRefExpr 0x305de40 <col:32> 'size_t':'unsigned long' lvalue ParmVar 0x305da50 'dim' 'size_t':'unsigned long'
|       |-UnaryOperator 0x305dee8 <col:37, col:39> 'size_t':'unsigned long' lvalue prefix '++'
|       | `-DeclRefExpr 0x305dec0 <col:39> 'size_t':'unsigned long' lvalue Var 0x305dbc0 'i' 'size_t':'unsigned long'
|       `-ForStmt 0x305e4c0 <line:79:13, line:80:44>
|         |-DeclStmt 0x305dfb0 <line:79:18, col:30>
|         | `-VarDecl 0x305df20 <col:18, col:29> col:25 used k 'size_t':'unsigned long' cinit
|         |   `-ImplicitCastExpr 0x305df98 <col:29> 'size_t':'unsigned long' <IntegralCast>
|         |     `-IntegerLiteral 0x305df78 <col:29> 'int' 0
|         |-<<<NULL>>>
|         |-BinaryOperator 0x305e048 <col:32, col:36> '_Bool' '<'
|         | |-ImplicitCastExpr 0x305e018 <col:32> 'size_t':'unsigned long' <LValueToRValue>
|         | | `-DeclRefExpr 0x305dfc8 <col:32> 'size_t':'unsigned long' lvalue Var 0x305df20 'k' 'size_t':'unsigned long'
|         | `-ImplicitCastExpr 0x305e030 <col:36> 'size_t':'unsigned long' <LValueToRValue>
```

```
                  `-DeclRefExpr 0x305dff0 <col:36> 'size_t':'unsigned long' lvalue ParmVar 0x305da50 'dim' 'size_t':'unsigned long'
          `-UnaryOperator 0x305e098 <col:41, col:43> 'size_t':'unsigned long' lvalue prefix '++'
            `-DeclRefExpr 0x305e070 <col:43> 'size_t':'unsigned long' lvalue Var 0x305dbc0 'i' 'size_t':'unsigned long'
  `-CompoundAssignOperator 0x305e488 <line:80:17, col:44> 'float' lvalue '+=' ComputeLHSTy='float' ComputeResultTy='float'
    |-ArraySubscriptExpr 0x305e1b8 <col:17, col:23> 'float' lvalue
    | |-ImplicitCastExpr 0x305e188 <col:17, col:20> 'float *' <LValueToRValue>
    | | `-ArraySubscriptExpr 0x305e138 <col:17, col:20> 'float *' lvalue
    | |   |-ImplicitCastExpr 0x305e108 <col:17> 'float **' <LValueToRValue>
    | |   | `-DeclRefExpr 0x305e0b8 <col:17> 'float **' lvalue ParmVar 0x305d9e0 'C' 'float **'
    | |   `-ImplicitCastExpr 0x305e120 <col:19> 'size_t':'unsigned long' <LValueToRValue>
    | |     `-DeclRefExpr 0x305e0e0 <col:19> 'size_t':'unsigned long' lvalue Var 0x305dbc0 'i' 'size_t':'unsigned long'
    | `-ImplicitCastExpr 0x305e1a0 <col:22> 'size_t':'unsigned long' <LValueToRValue>
    |   `-DeclRefExpr 0x305e160 <col:22> 'size_t':'unsigned long' lvalue Var 0x305dd70 'j' 'size_t':'unsigned long'
    `-BinaryOperator 0x305e460 <col:28, col:44> 'float' '*'
      |-ImplicitCastExpr 0x305e430 <col:28, col:34> 'float' <LValueToRValue>
      | `-ArraySubscriptExpr 0x305e2e0 <col:28, col:34> 'float' lvalue
      |   |-ImplicitCastExpr 0x305e2b0 <col:28, col:31> 'float *' <LValueToRValue>
      |   | `-ArraySubscriptExpr 0x305e260 <col:28, col:31> 'float *' lvalue
      |   |   |-ImplicitCastExpr 0x305e230 <col:28> 'float **' <LValueToRValue>
      |   |   | `-DeclRefExpr 0x305e1e0 <col:28> 'float **' lvalue ParmVar 0x305d8e0 'A' 'float **'
      |   |   `-ImplicitCastExpr 0x305e248 <col:30> 'size_t':'unsigned long' <LValueToRValue>
      |   |     `-DeclRefExpr 0x305e208 <col:30> 'size_t':'unsigned long' lvalue Var 0x305dbc0 'i' 'size_t':'unsigned long'
      |   `-ImplicitCastExpr 0x305e2c8 <col:33> 'size_t':'unsigned long' <LValueToRValue>
      |     `-DeclRefExpr 0x305e288 <col:33> 'size_t':'unsigned long' lvalue Var 0x305df20 'k' 'size_t':'unsigned long'
      `-ImplicitCastExpr 0x305e448 <col:38, col:44> 'float' <LValueToRValue>
        `-ArraySubscriptExpr 0x305e408 <col:38, col:44> 'float' lvalue
          |-ImplicitCastExpr 0x305e3d8 <col:38, col:41> 'float *' <LValueToRValue>
          | `-ArraySubscriptExpr 0x305e388 <col:38, col:41> 'float *' lvalue
          |   |-ImplicitCastExpr 0x305e358 <col:38> 'float **' <LValueToRValue>
          |   | `-DeclRefExpr 0x305e308 <col:38> 'float **' lvalue ParmVar 0x305d970 'B' 'float **'
          |   `-ImplicitCastExpr 0x305e370 <col:40> 'size_t':'unsigned long' <LValueToRValue>
          |     `-DeclRefExpr 0x305e330 <col:40> 'size_t':'unsigned long' lvalue Var 0x305df20 'k' 'size_t':'unsigned long'
          `-ImplicitCastExpr 0x305e3f0 <col:43> 'size_t':'unsigned long' <LValueToRValue>
            `-DeclRefExpr 0x305e3b0 <col:43> 'size_t':'unsigned long' lvalue Var 0x305dd70 'j' 'size_t':'unsigned long'
```

References

[1] "Cuda," 2016, license: http://docs.nvidia/com/cuda/eula. [Online]. Available: http://www.nvidia.com/object/cuda_home_new.html

[2] "Oclint," 2016, license: Modified BSD. [Online]. Available: http://oclint.org/

[3] "Eigen," 2016, license: MPL2. [Online]. Available: http://eigen.tuxfamily.org

[4] "Thrust," 2016, license: Apache 2.0. [Online]. Available: http://thrust.github.io

[5] "Cublas," 2016, license: http://docs.nvidia.com/cuda/eula/. [Online]. Available: https://developer.nvidia.com/cublas

[6] R. Espasa and M. Valero, "Exploiting instruction-and data-level parallelism," *IEEE micro*, vol. 17, no. 5, pp. 20–27, 1997.

[7] "Auto-vectorization in llvm." [Online]. Available: http://llvm.org/docs/Vectorizers.html

[8] Y. Sui, X. Fan, H. Zhou, and J. Xue, "Loop-oriented array-and field-sensitive pointer analysis for automatic simd vectorization," in *Proceedings of the 17th ACM SIGPLAN/SIGBED Conference on Languages, Compilers, Tools, and Theory for Embedded Systems.* ACM, 2016, pp. 41–51.

[9] M. Papadrakakis, G. Stavroulakis, and A. Karatarakis, "A new era in scientific computing: Domain decomposition methods in hybrid cpu–gpu architectures," *Computer Methods in Applied Mechanics and Engineering*, vol. 200, no. 13, pp. 1490–1508, 2011.

[10] K. Nagarajan, B. Holland, A. D. George, K. C. Slatton, and H. Lam, "Accelerating machine-learning algorithms on fpgas using pattern-based decomposition," *Journal of Signal Processing Systems*, vol. 62, no. 1, pp. 43–63, 2011.

[11] S. Sengupta, M. Harris, Y. Zhang, and J. D. Owens, "Scan primitives for gpu computing," in *Graphics hardware*, vol. 2007, 2007, pp. 97–106.

[12] A. Bessey, K. Block, B. Chelf, A. Chou, B. Fulton, S. Hallem, C. Henri-Gros, A. Kamsky, S. McPeak, and D. Engler, "A few billion lines of code later: using static analysis to find bugs in the real world," *Communications of the ACM*, vol. 53, no. 2, pp. 66–75, 2010.

[13] T. Ball and S. K. Rajamani, "Automatically validating temporal safety properties of interfaces," in *Proceedings of the 8th international SPIN workshop on Model checking of software.* Springer-Verlag New York, Inc., 2001, pp. 103–122.

[14] W. R. Bush, J. D. Pincus, and D. J. Sielaff, "A static analyzer for finding dynamic programming errors," *Software-Practice and Experience*, vol. 30, no. 7, pp. 775–802, 2000.

[15] C. Woolley, "High-productivity cuda programming," 2013. [Online]. Available: http://on-demand.gputechconf.com/gtc/2013/presentations/ S3008-High-Productivity-CUDA-Programming.pdf

[16] G. S. Almasi, *Highly Parallel Processing (The Benjamin/Cummings series in computer science and engineering).* Benjamin-Cummings Publishing Co.,Subs. of Addison Wesley Longman,US, 1987. [Online]. Available: http: //www.amazon.com/Parallel-Processing-Benjamin-Cummings-engineering/ dp/0805301771%3FSubscriptionId%3D0JYN1NVW651KCA56C102% 26tag%3Dtechkie-20%26linkCode%3Dxm2%26camp%3D2025%26creative% 3D165953%26creativeASIN%3D0805301771

[17] M. Harris, "Parallel prefix sum (scan) with cuda," 2007. [Online]. Available: http://www.eecs.umich.edu/courses/eecs570/hw/parprefix.pdf

[18] O. Inzunza-Monreal, "Performance of parallel processing on processing units."

[19] "Clang features and goals," 2016. [Online]. Available: http://clang.llvm.org/ features.html

[20] "Json for modern c++," 2016, license: MIT. [Online]. Available: https://nlohmann.github.io/json/

[21] "hayai," 2016, license: https://github.com/nickbruun/hayai/blob/master/LICENSE.md. [Online]. Available: https://github.com/nickbruun/hayai